

GPGPU

A look into OpenCL

Author : Frédéric Dubouchet

Professor : Paul Albuquerque

Plan

- What is OpenCL?
- Practical simple example
- Advance examples in OpenCL
- Conclusion
- Questions

What is OpenCL?

Open Computing Language

Language and API

OpenCL Objects

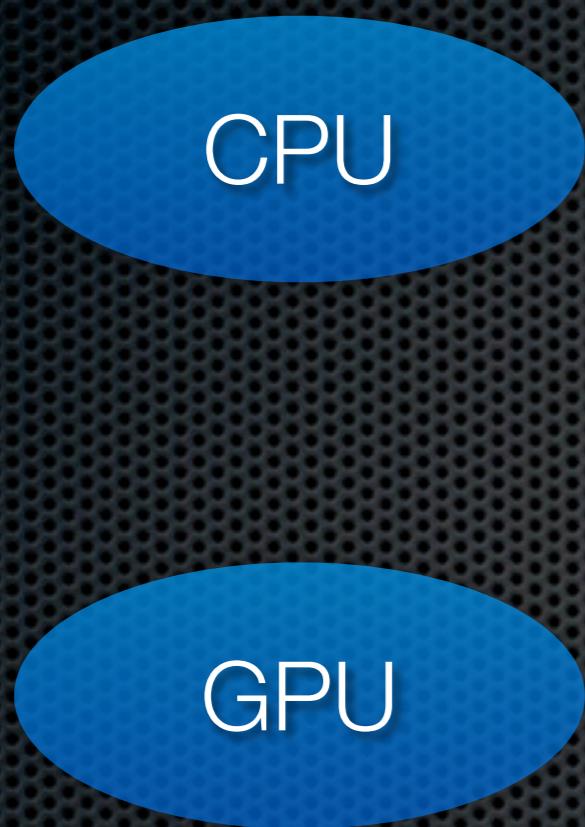
OpenCL Device Model

OpenCL C++ Interface

Open Computing Language

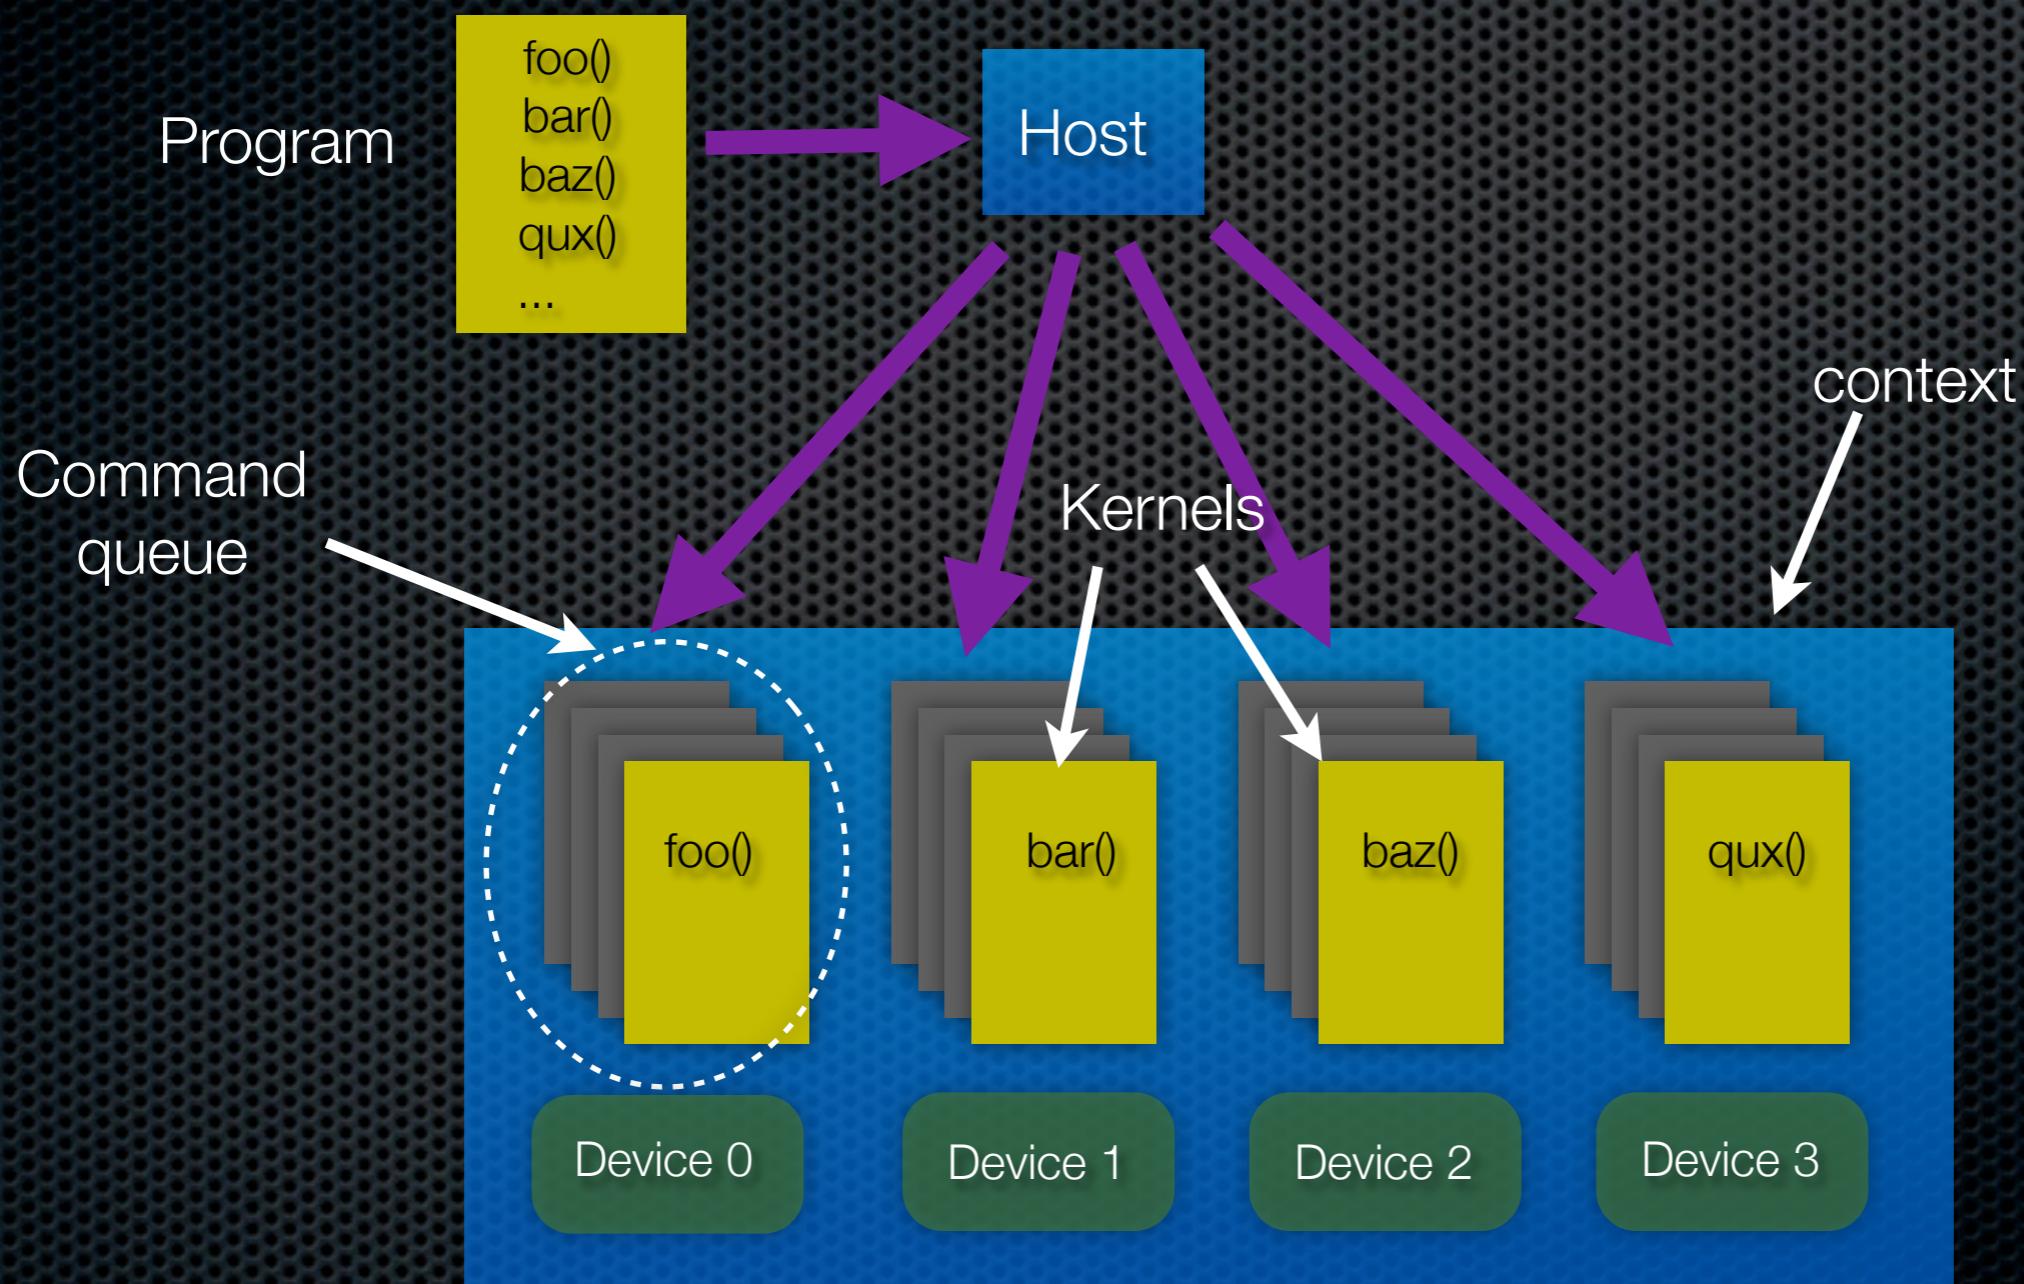
- General Purpose Computing Language!
 - Parallel computing on many platform
 - computer mobile web
- GPGPU but not only!
 - CPU / GPU / Specialized Hardware
- Maintain by the Khronos group (OpenGL)
 - Apple, Nvidia, AMD, Intel and many others...
- Open! not another proprietary interface

Language and API

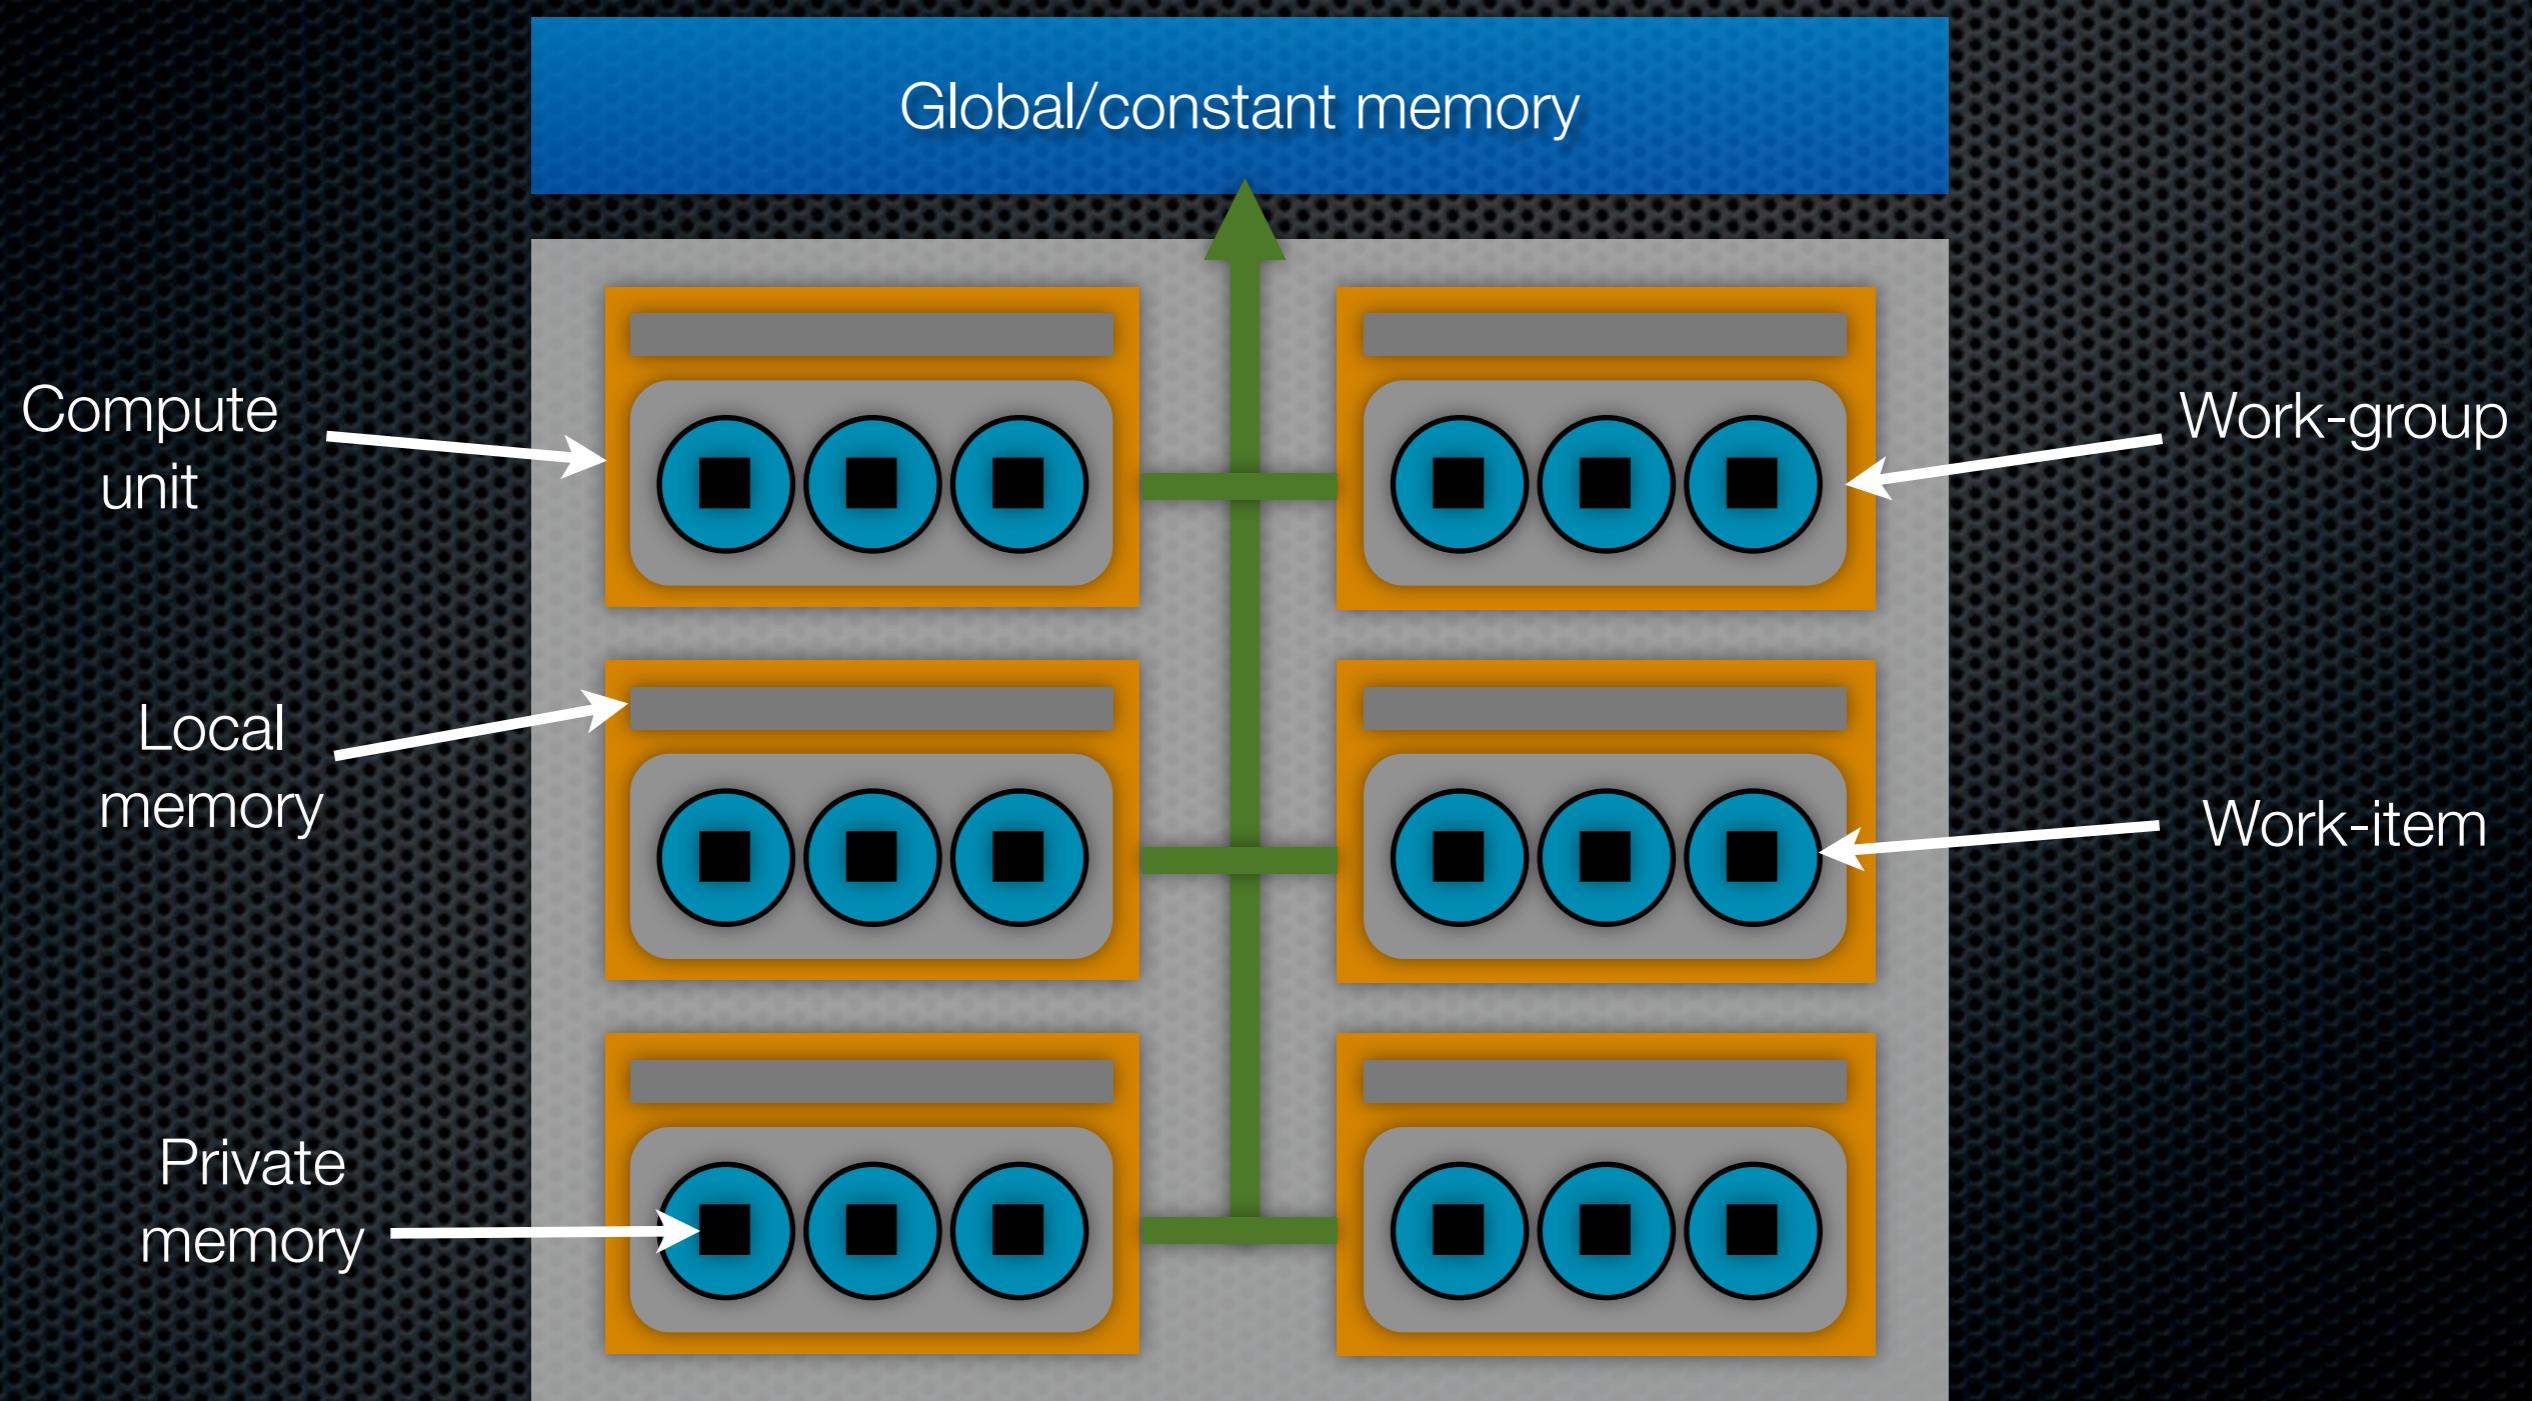


- An API (run on the Host)
 - in C but with interface to many other languages
- A language (run on the Device)
 - vector oriented
 - C99 inspired

OpenCL Objects

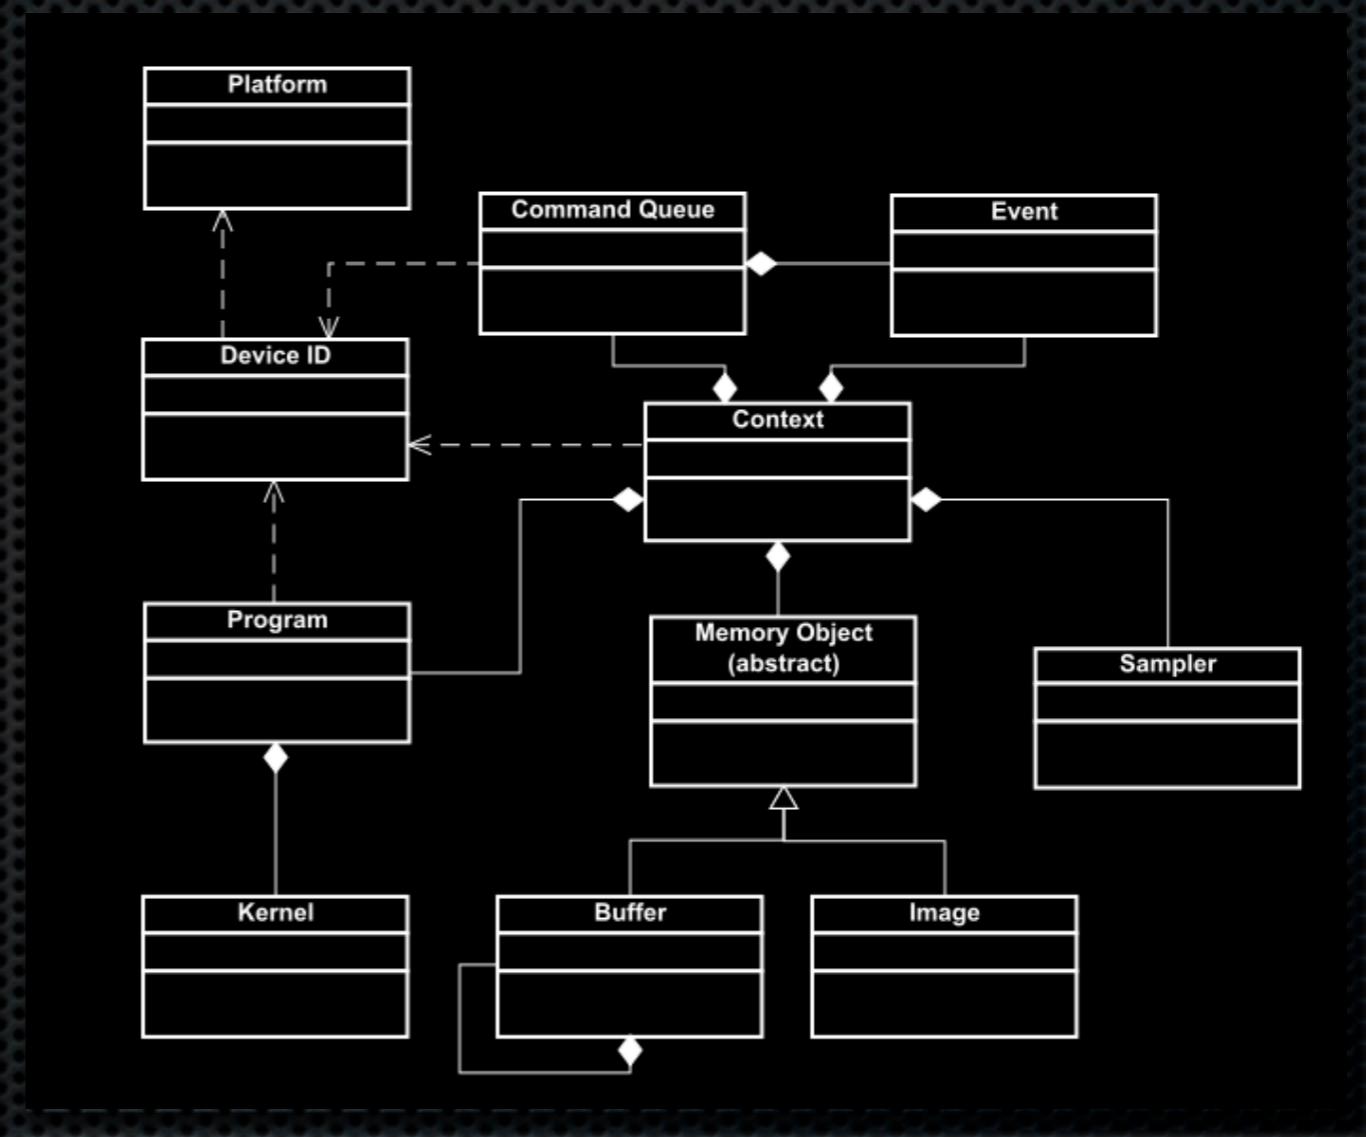


OpenCL Device Model



OpenCL C++ interface

- Officially published by Khronos
- Template based
- Header only
- 100% portable
- Object Oriented



Practical Simple Example

Host code

Device code

Host code (1/2)

- ❖ Headers include
- ❖ Get Platform
- ❖ Get Device
- ❖ Create Context
- ❖ Create Command Queue
- ❖ Load & Build
- ❖ Select Kernel

```
// simple OpenCL example

#include <iostream>
#include <fstream>
#include <vector>
#define __CL_ENABLE_EXCEPTIONS
#include <CL/cl.hpp>

const unsigned int vector_size = 1024;

int main(int ac, char** av) {
    try {
        // get the device (here the GPU)
        std::vector<cl::Platform> platforms;
        std::vector<cl::Device> devices_;
        cl::Platform::get(&platforms);
        platforms[0].getDevices(CL_DEVICE_TYPE_GPU, &devices_);
        cl_context_properties properties[] = {
            CL_CONTEXT_PLATFORM,
            (cl_context_properties)(platforms[0])(),
            0
        };
        cl::Context context_ = cl::Context(CL_DEVICE_TYPE_GPU, properties);
        devices_ = context_.getInfo<CL_CONTEXT_DEVICES>();
        cl::CommandQueue queue_(context_, devices_[0]);
        // load the kernel code
        std::ifstream ifs("./simple.cl");
        std::string str(
            (std::istreambuf_iterator<char>(ifs)),
            std::istreambuf_iterator<char>());
        ifs.close();
        // compile
        cl::Program::Sources source(1, std::make_pair(str.c_str(), str.size()));
        cl::Program program_(context_, source);
        program_.build(devices_);
        // create the kernel
        cl::Kernel kernel_(program_, "simple");
```

Host code (2/2)

- Create and fill up Host buffers
- Create Device Buffers (2 in and 1 out)
- Set Kernel Arguments
- Execute the Kernel
- Read the buffers
- Exception Handling

```
// prepare the buffers
std::vector<float> in1(vector_size);
std::vector<float> in2(vector_size);
for (std::vector<float>::iterator ite = in1.begin(); ite != in1.end(); ++ite)
    (*ite) = (float)random();
for (std::vector<float>::iterator ite = in2.begin(); ite != in2.end(); ++ite)
    (*ite) = (float)random();
cl::Buffer buf_in1_ = cl::Buffer(
    context_,
    CL_MEM_READ_ONLY | CL_MEM_COPY_HOST_PTR,
    sizeof(float) * vector_size,
    (void*)&in1);
cl::Buffer buf_in2_ = cl::Buffer(
    context_,
    CL_MEM_READ_ONLY | CL_MEM_COPY_HOST_PTR,
    sizeof(float) * vector_size,
    (void*)&in2);
cl::Buffer buf_out_ = cl::Buffer(
    context_,
    CL_MEM_WRITE_ONLY,
    sizeof(float) * vector_size);
// set the arguments
kernel_.setArg(0, buf_in1_);
kernel_.setArg(1, buf_in2_);
kernel_.setArg(2, buf_out_);
// wait to the command queue to finish before proceeding
queue_.finish();
// run the kernel
std::vector<float> out(vector_size);
queue_.enqueueNDRangeKernel(
    kernel_,
    cl::NullRange,
    cl::NDRange(vector_size),
    cl::NullRange);
queue_.finish();
// get the result out
queue_.enqueueReadBuffer(
    buf_out_,
    CL_TRUE,
    0,
    vector_size * sizeof(float),
    &out[0]);
queue_.finish();
std::cout << "Operation successfull" << std::endl;
} catch (cl::Error& er) {
    std::cerr << "Exception(CL) : " << er.what() << std::endl;
} catch (std::exception& ex) {
    std::cerr << "Exception(STL) : " << ex.what() << std::endl;
}
return 0;
```

Device Code

- Kernel simple
- global in buffers
(in1 & in2)
- global out buffer
- get the work item ID
- compute

```
// very simple kernel

__kernel void simple(
    __global read_only float* in1,
    __global read_only float* in2,
    __global write_only float* out)
{
    const uint pos = get_global_id(0);
    out[pos] = in1[pos] * in2[pos];
}
```

Advance examples in OpenCL

Julia Set

Cellular automata

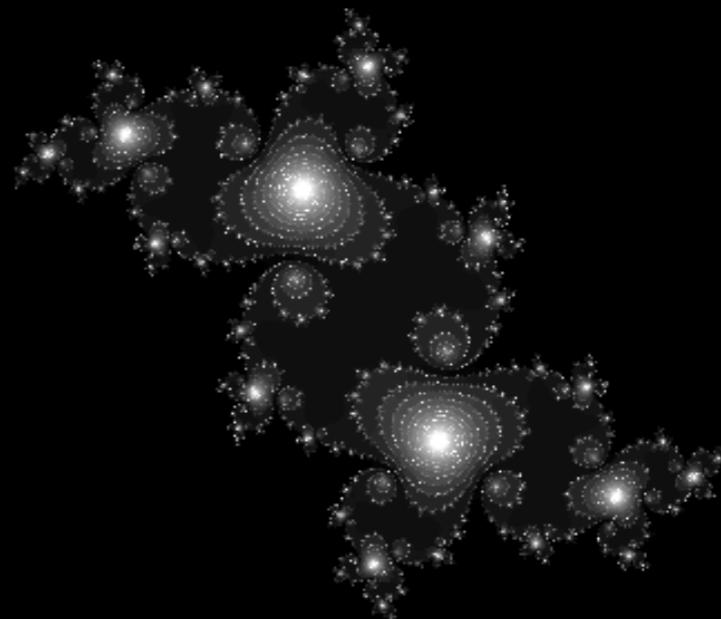
Floyd-Warshall

Fast Fourier Transform

Julia Set

Kernel view

Performances



Kernel View

- Utility complex operators
- Get WorkItem position (d)
- Get grid size (m)
- Compute the position in C
- Compute Julia
- Set the output buffer

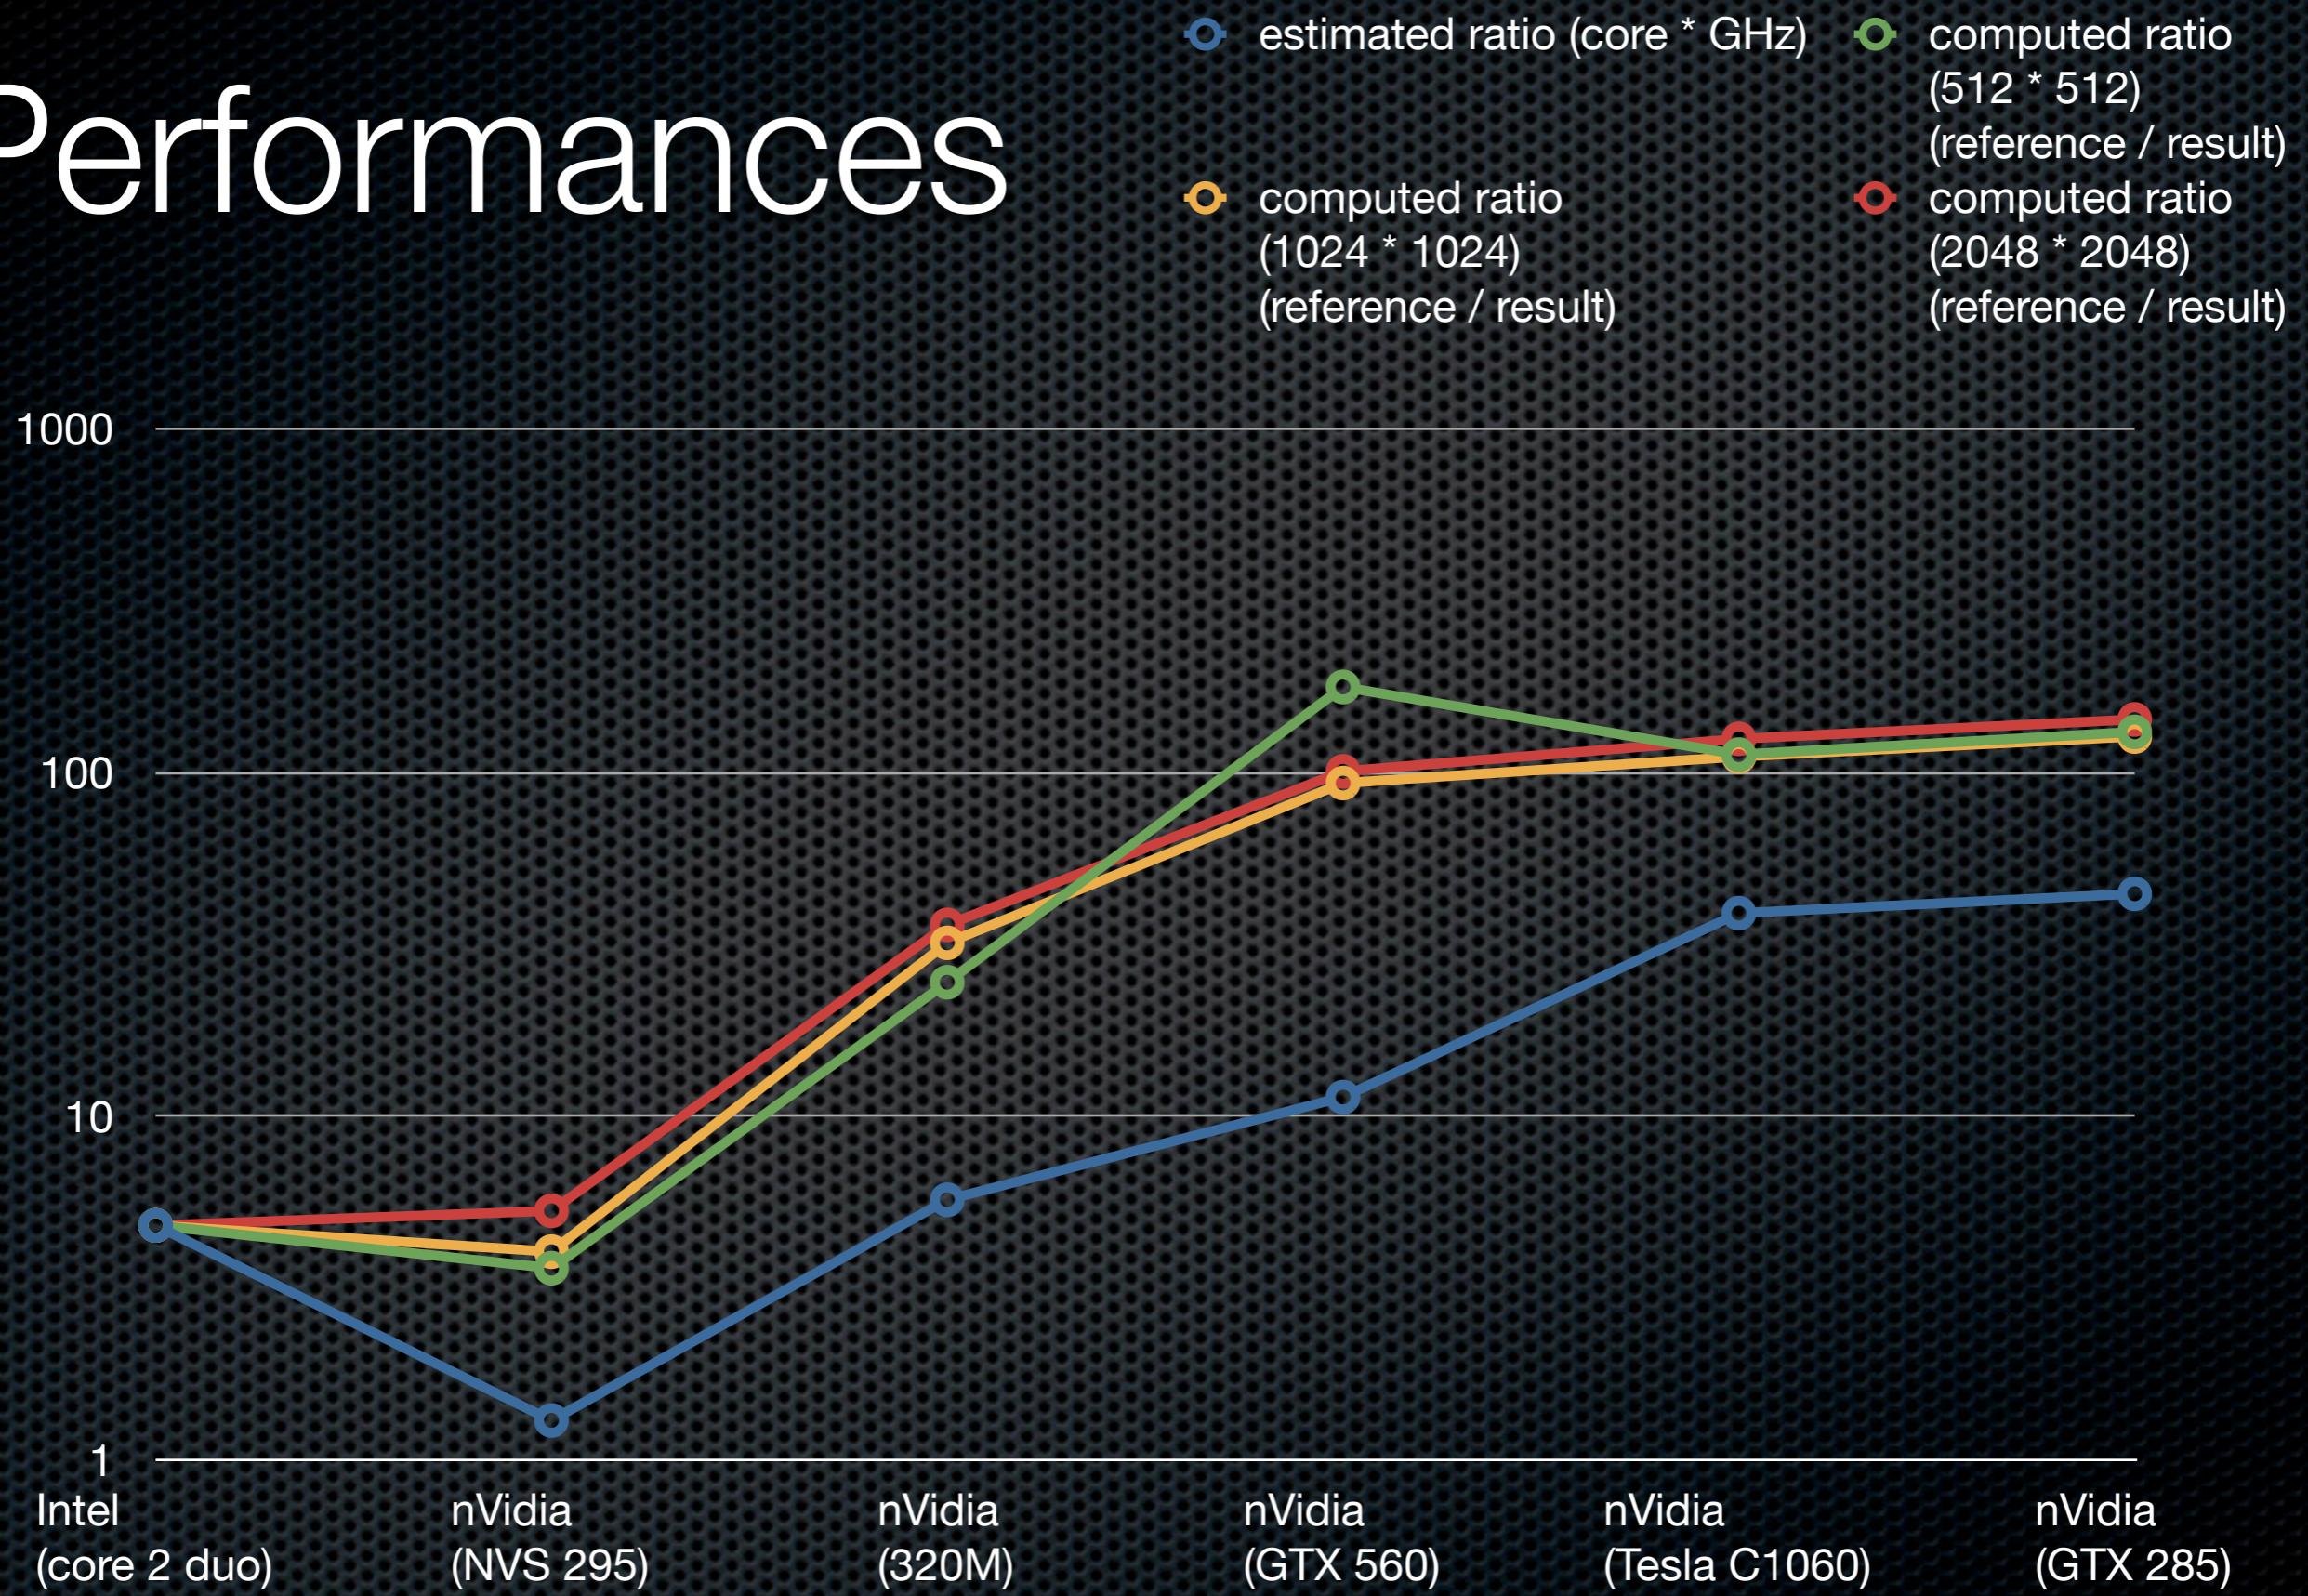
```
float2 square_c(float2 c) {
    return (float2)(c.x * c.x - c.y * c.y, 2.0f * c.y * c.x);
}

float2 add_c(float2 c1, float2 c2) {
    return (float2)(c1.x + c2.x, c1.y + c2.y);
}

float square_norm_c(float2 c) {
    return c.x * c.x + c.y * c.y;
}

__kernel void julia_buffer(
    __global float* out,
    float2 c,
    uint max_iteration)
{
    const uint2 d = (uint2)(get_global_id(0), get_global_id(1));
    const uint2 m = (uint2)(get_global_size(0), get_global_size(1));
    float2 pos = (float2)(
        (((float)d.x / (float)m.x) * 4.0f) - 2.0f,
        (((float)d.y / (float)m.y) * 4.0f) - 2.0f);
    uint ite = 0;
    while ((square_norm_c(pos) < 4.0f) && (ite < max_iteration)) {
        pos = add_c(square_c(pos), c);
        ite++;
    }
    float val = (float)ite / (float)max_iteration;
    out[d.y * m.x + d.x] = val;
}
```

Performances

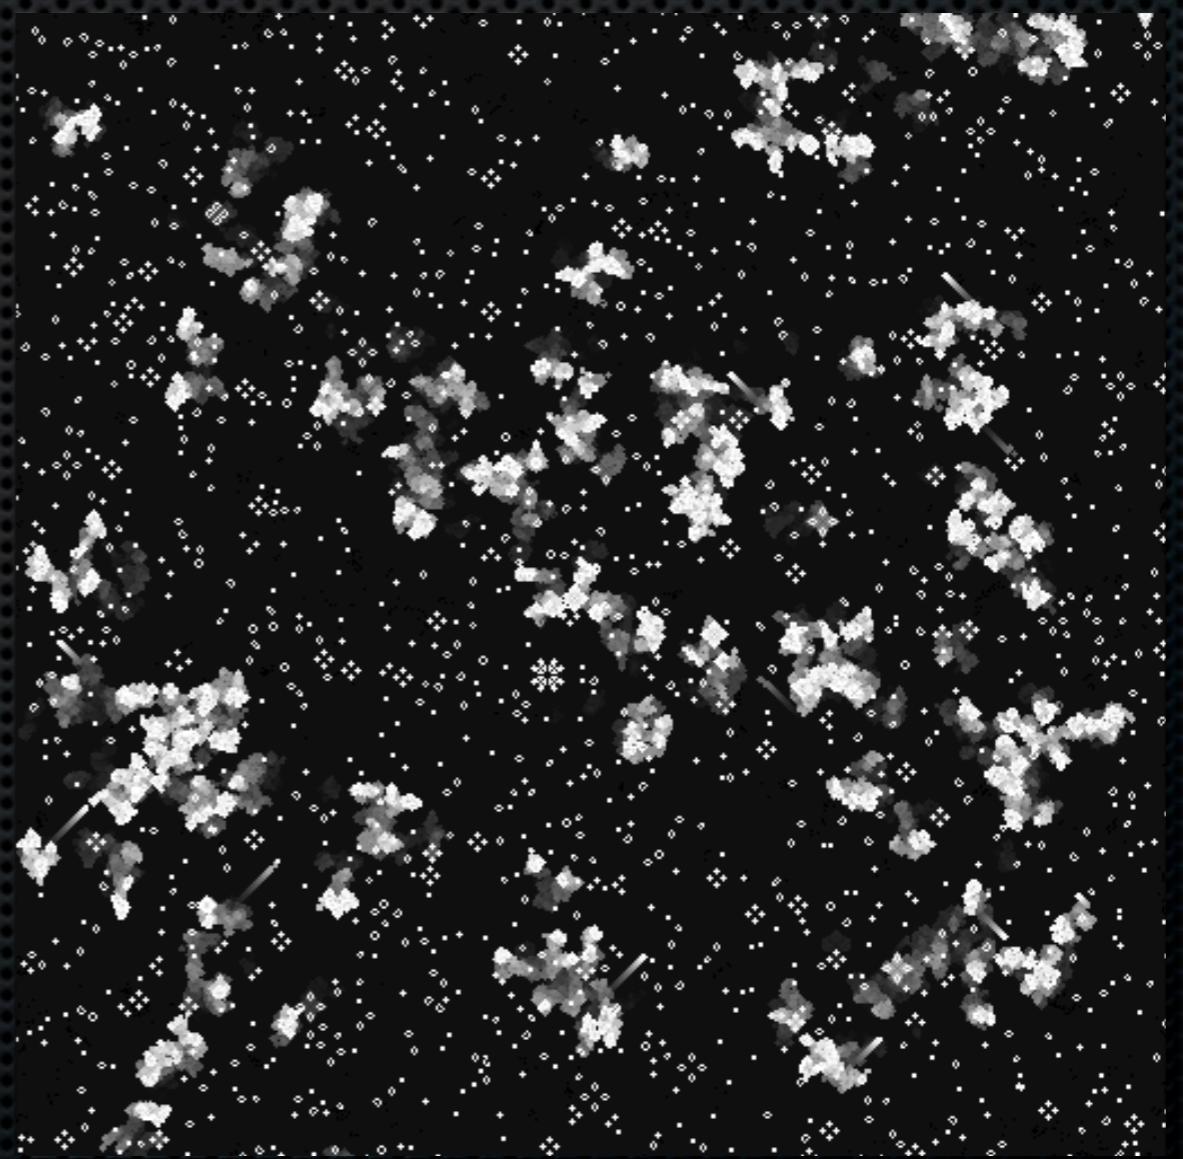


Cellular Automata

Kernel view

Performances

Against CUDA



Kernel View

- Utility complex operators
- Get WorkItem position (d)
- Get grid size (m)
- accumulate the value of neighbor
- Compute the result
 - using ternary operator

```
int pos2linear(
    const int2 position,
    const int2 m)
{
    return position.y * m.x + position.x;
}

int add(__global const float* in, const int2 position, const int2 m) {
    // empty border
    return
        ((position.x < 0 || position.x >= m.x) || (position.y < 0 || position.y >= m.y))
            ? 0.0f
            : (in[pos2linear(position, m)]) ? 1.0f : 0.0f;
}

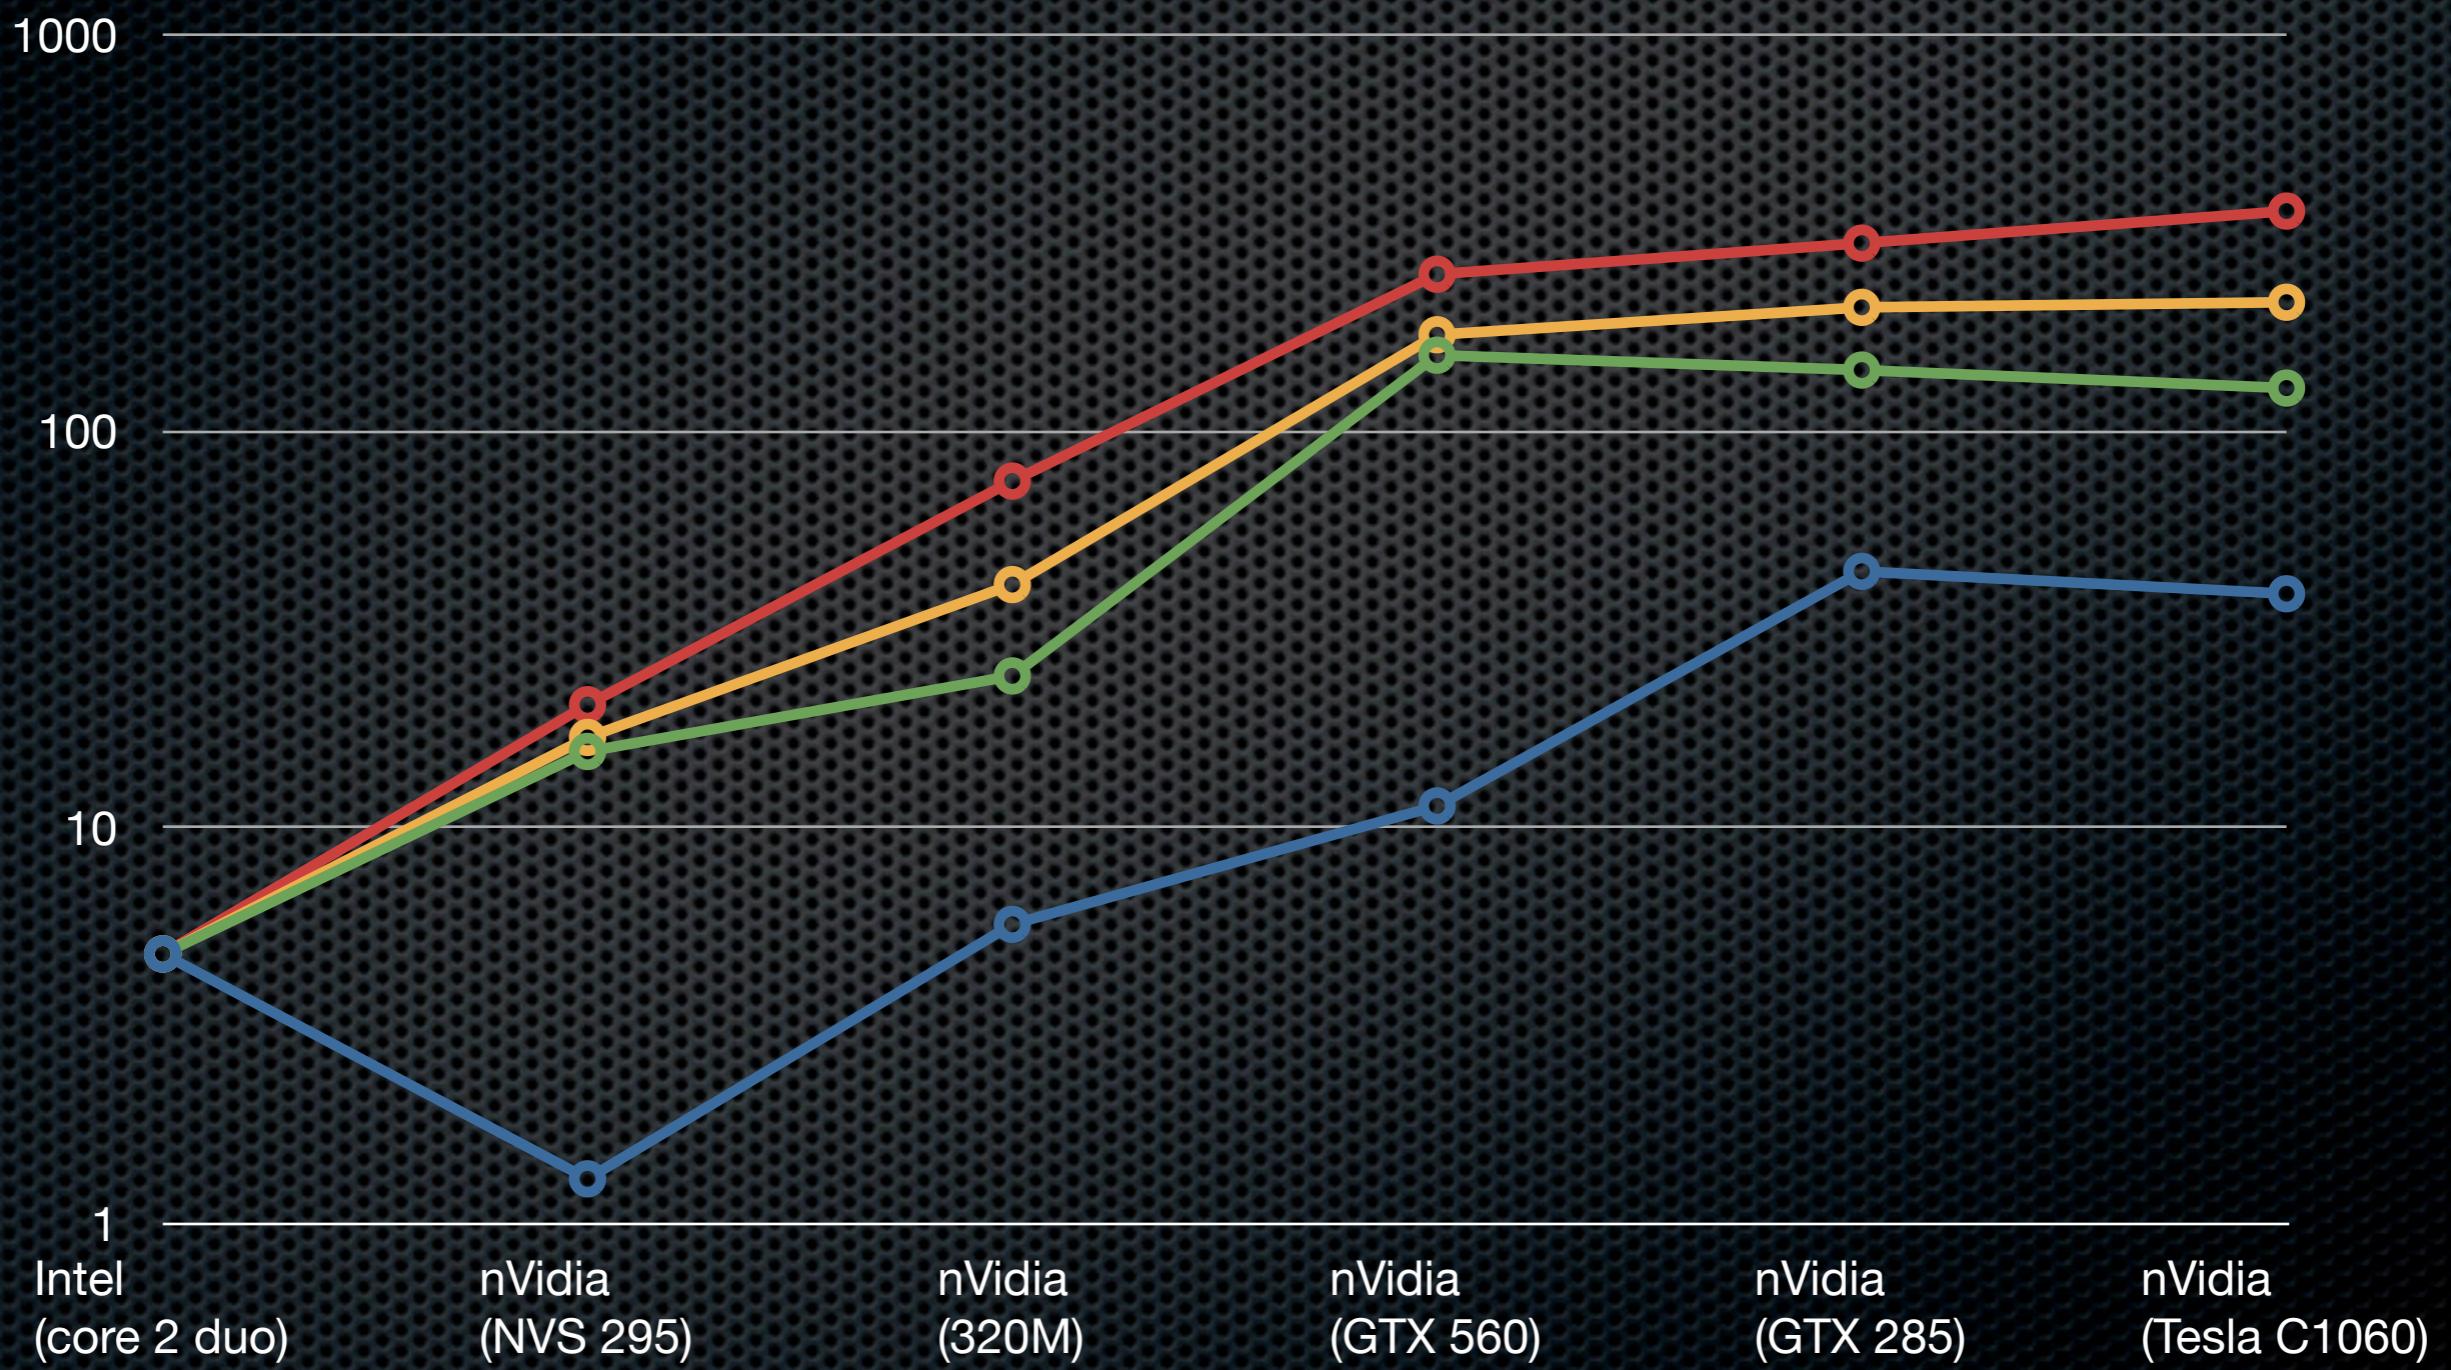
__kernel void life_buffer(
    __global float* in,
    __global float* out,
    unsigned int max_iterations)
{
    const int2 d = (int2)(get_global_id(0), get_global_id(1));
    const int2 m = (int2)(get_global_size(0), get_global_size(1));

    int position = pos2linear(d, m);
    int sum = 0;
    sum += add(in, (int2)(d.x - 1, d.y + 1), m);
    sum += add(in, (int2)(d.x      , d.y + 1), m);
    sum += add(in, (int2)(d.x + 1, d.y + 1), m);
    sum += add(in, (int2)(d.x - 1, d.y      ), m);
    sum += add(in, (int2)(d.x + 1, d.y      ), m);
    sum += add(in, (int2)(d.x - 1, d.y - 1), m);
    sum += add(in, (int2)(d.x      , d.y - 1), m);
    sum += add(in, (int2)(d.x + 1, d.y - 1), m);

    out[position] = ((sum == 3) || (in[position] && (sum == 2))) ? 1.0f : 0.0f;
}
```

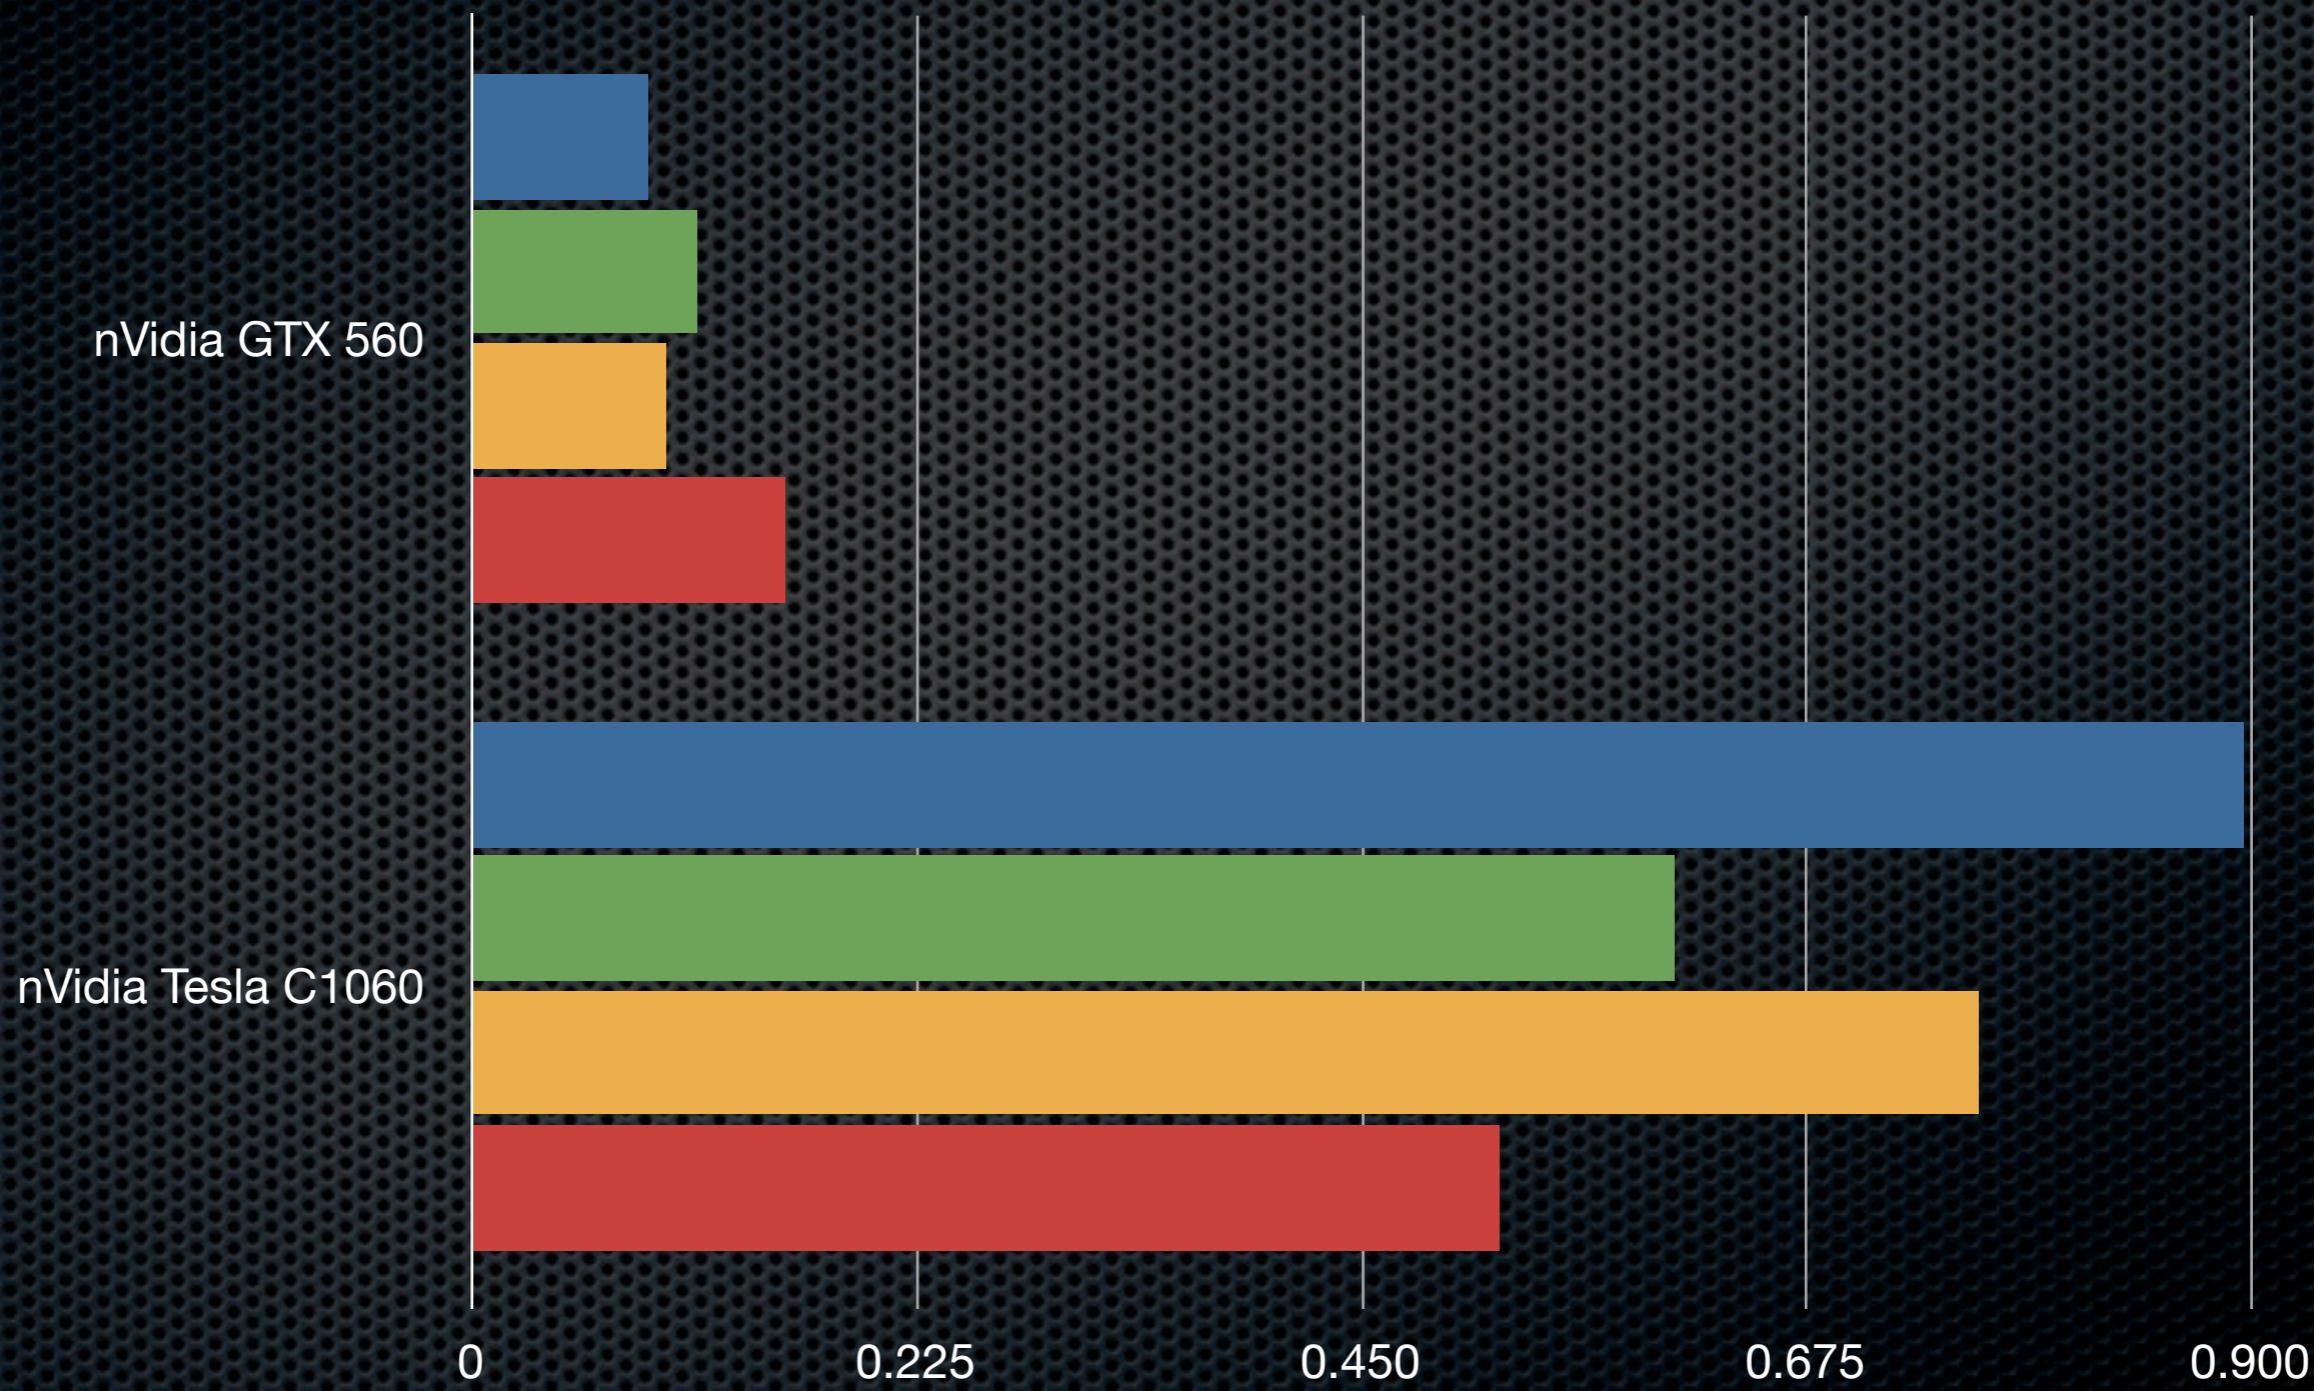
Performances

- estimated ratio (core * GHz)
- computed ratio (512 * 512) (reference / result)
- computed ratio (1024 * 1024) (reference / result)
- computed ratio (2048 * 2048) (reference / result)



Against CUDA

- CUDA (gpu)
- CUDA (gpushared)
- OpenCL (cl_buffer)
- OpenCL (cl_image2d_t)



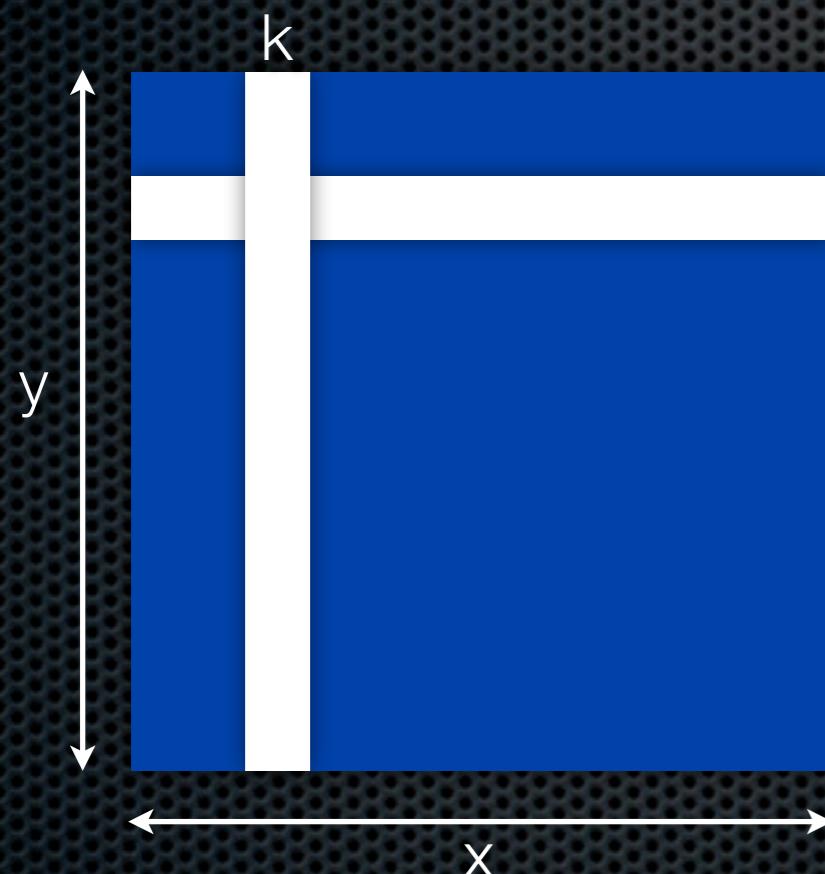
Floyd-Warshall

Method
Performances

Method

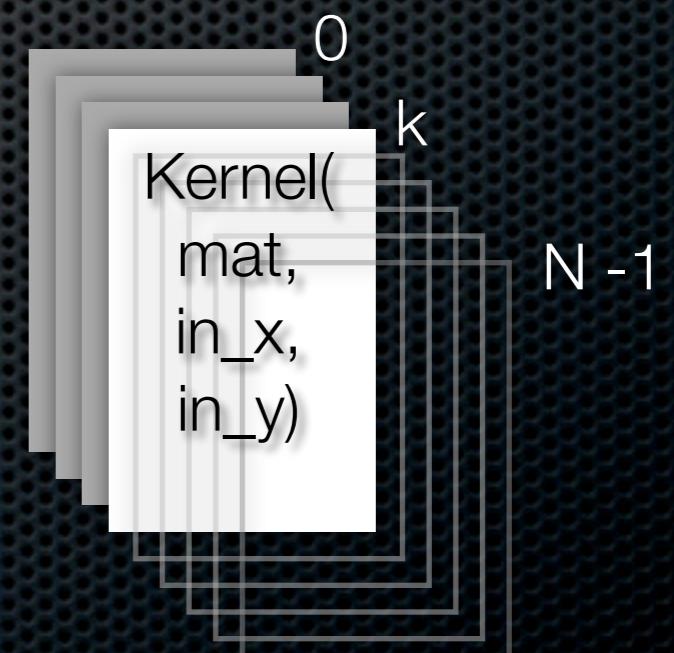
For each values of k (in $0..N-1$)

The Matrix

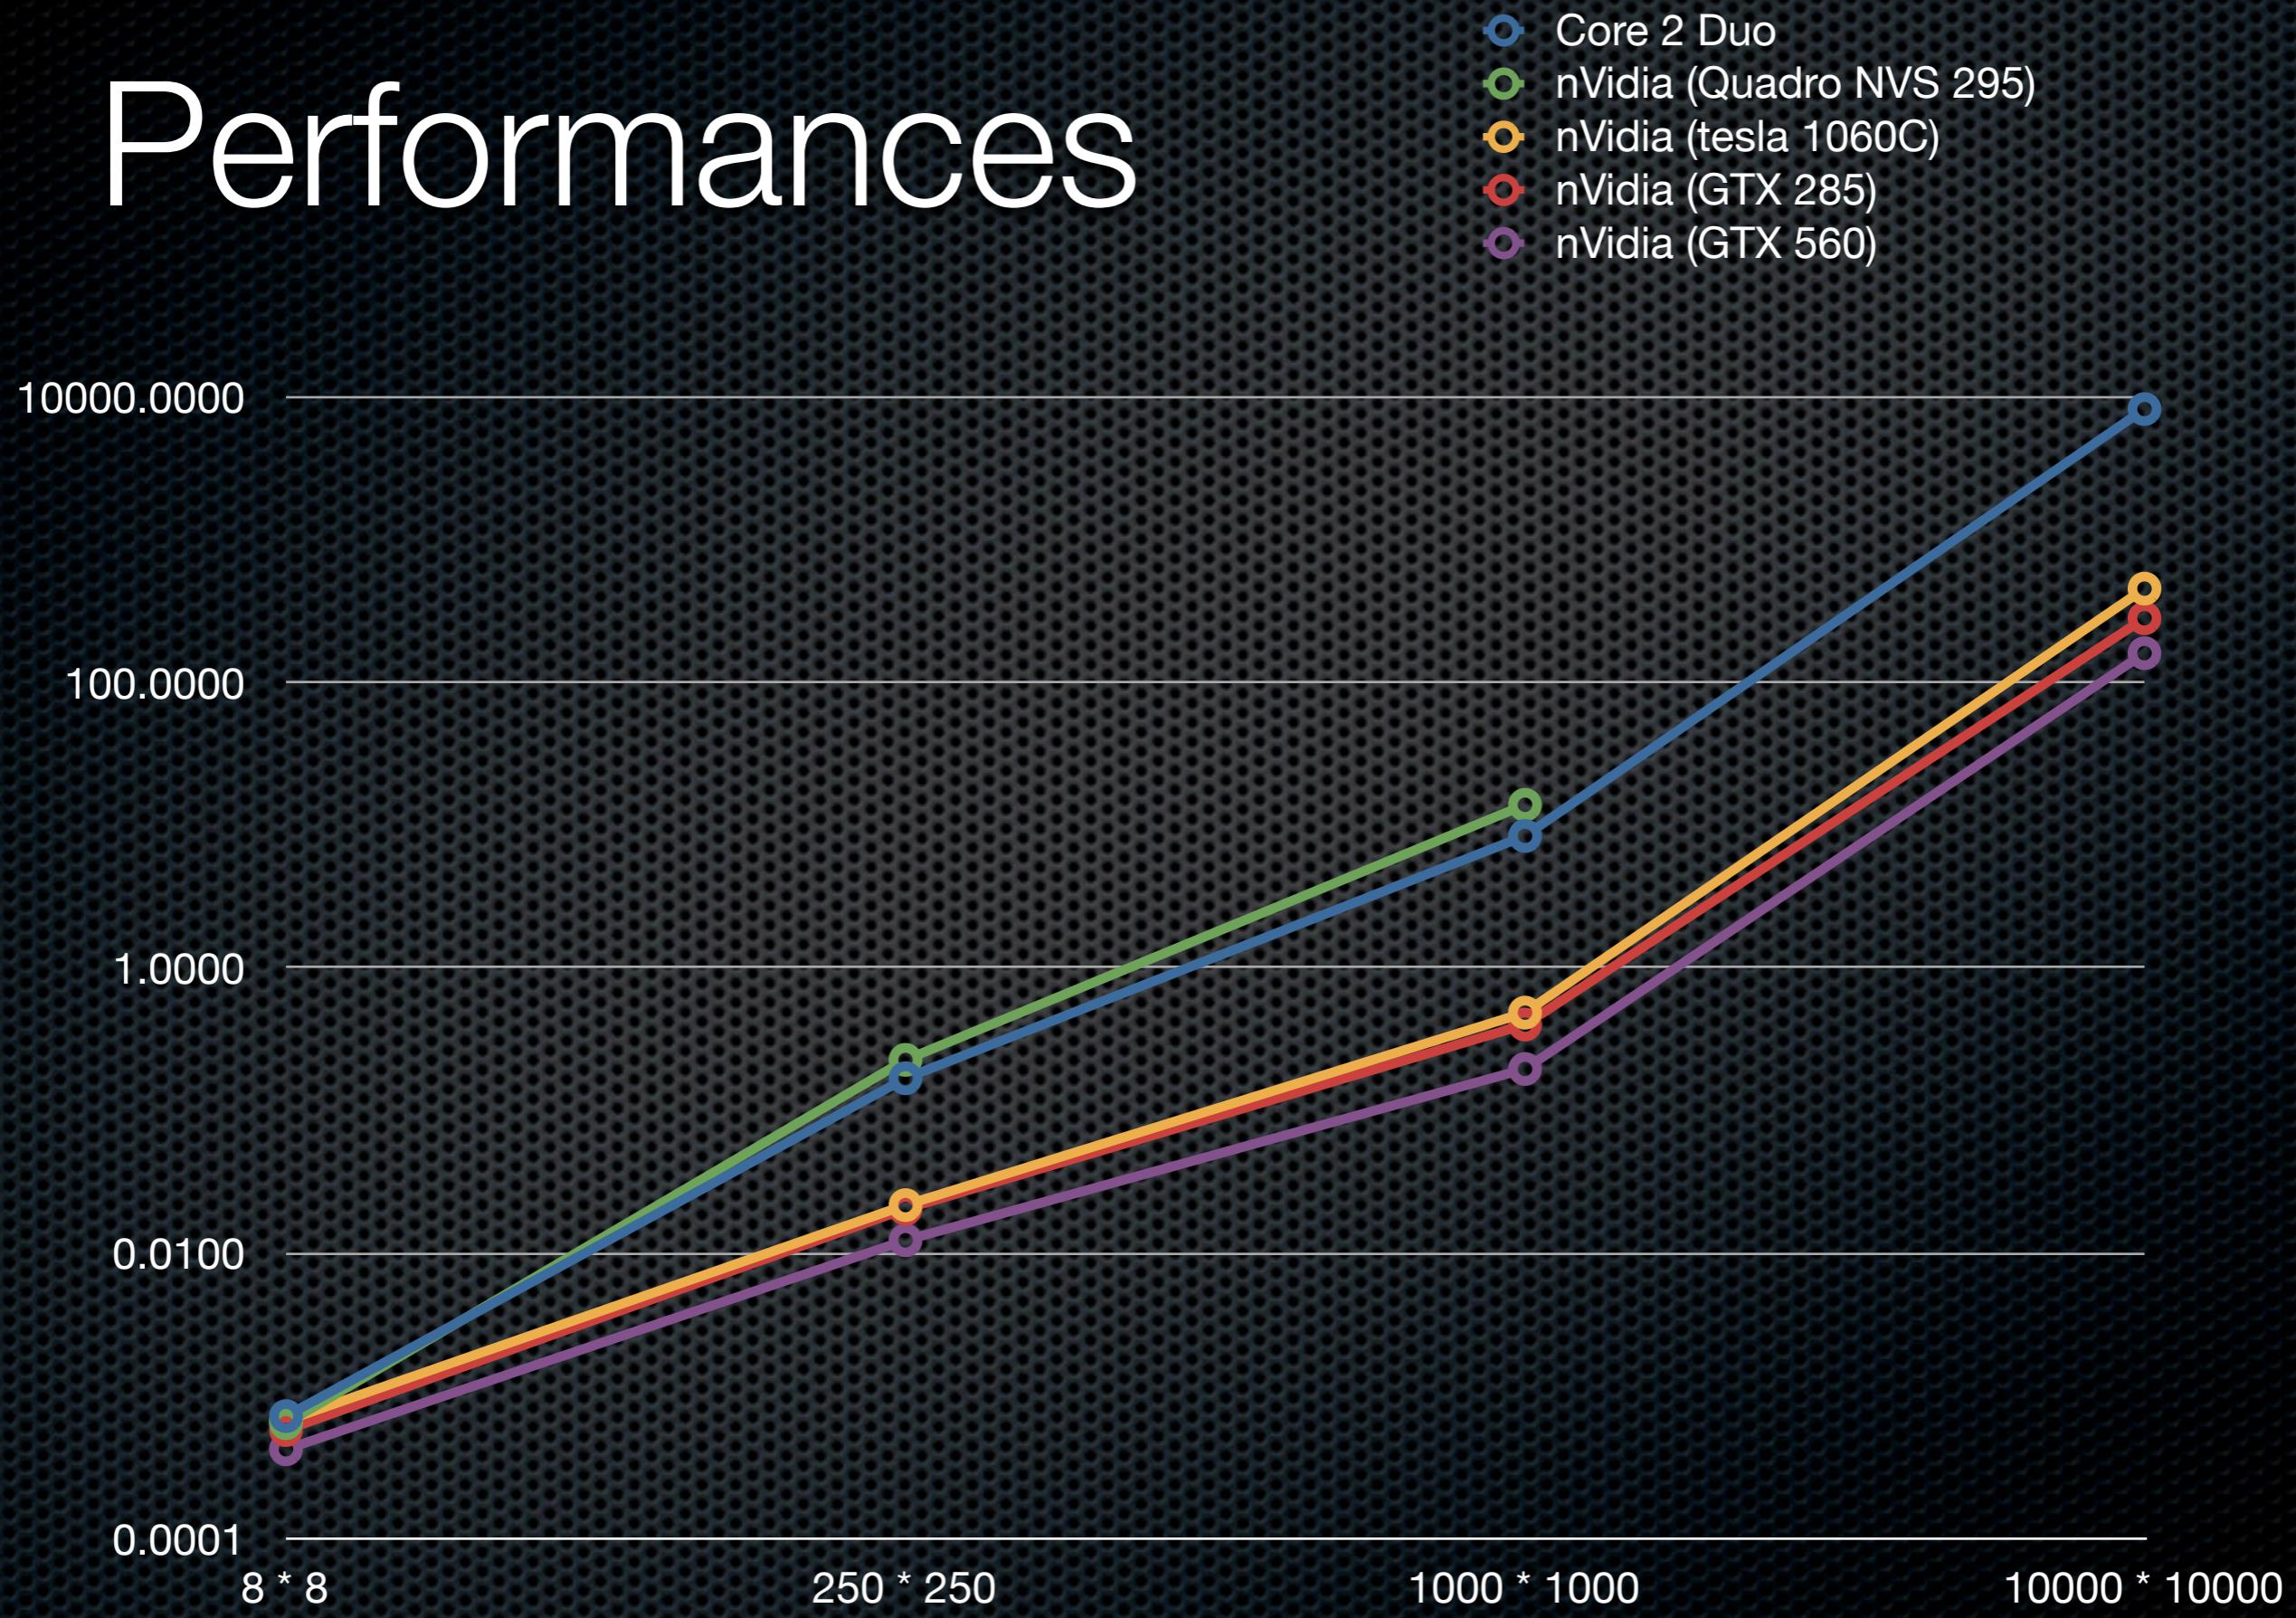


The Kernel Stack

mat (whole)
 $in_x = y[k]$
 $in_y = x[k]$



Performances



Fast Fourier Transform

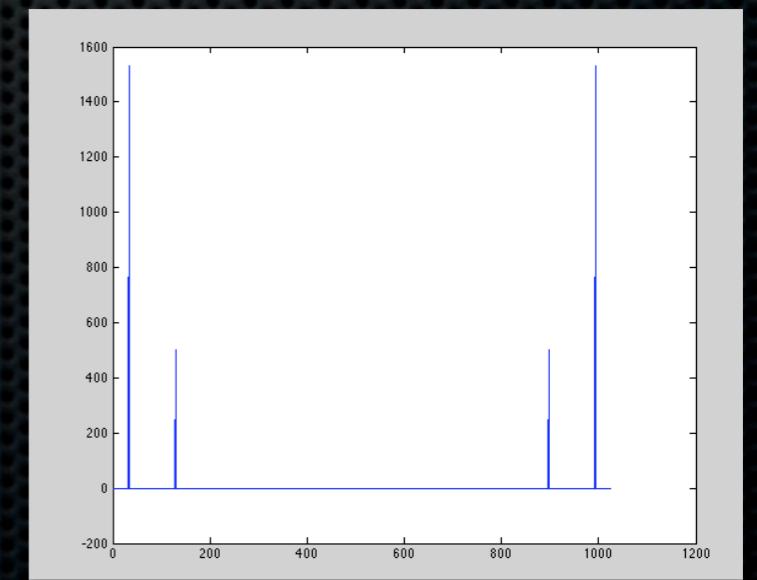
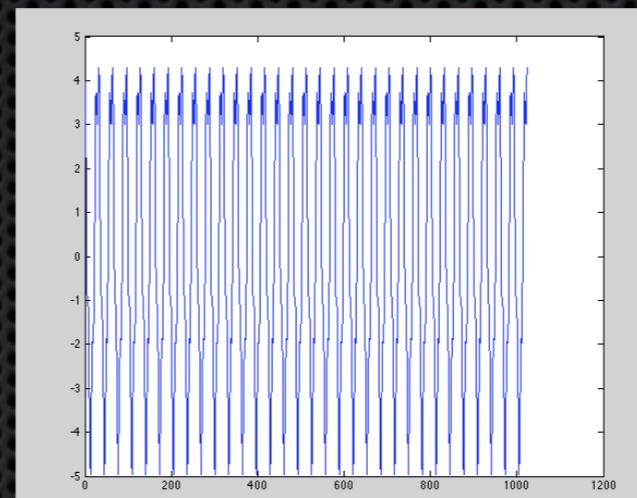
Fast Introduction

Kernel & FFTW

Performances

Fast Introduction

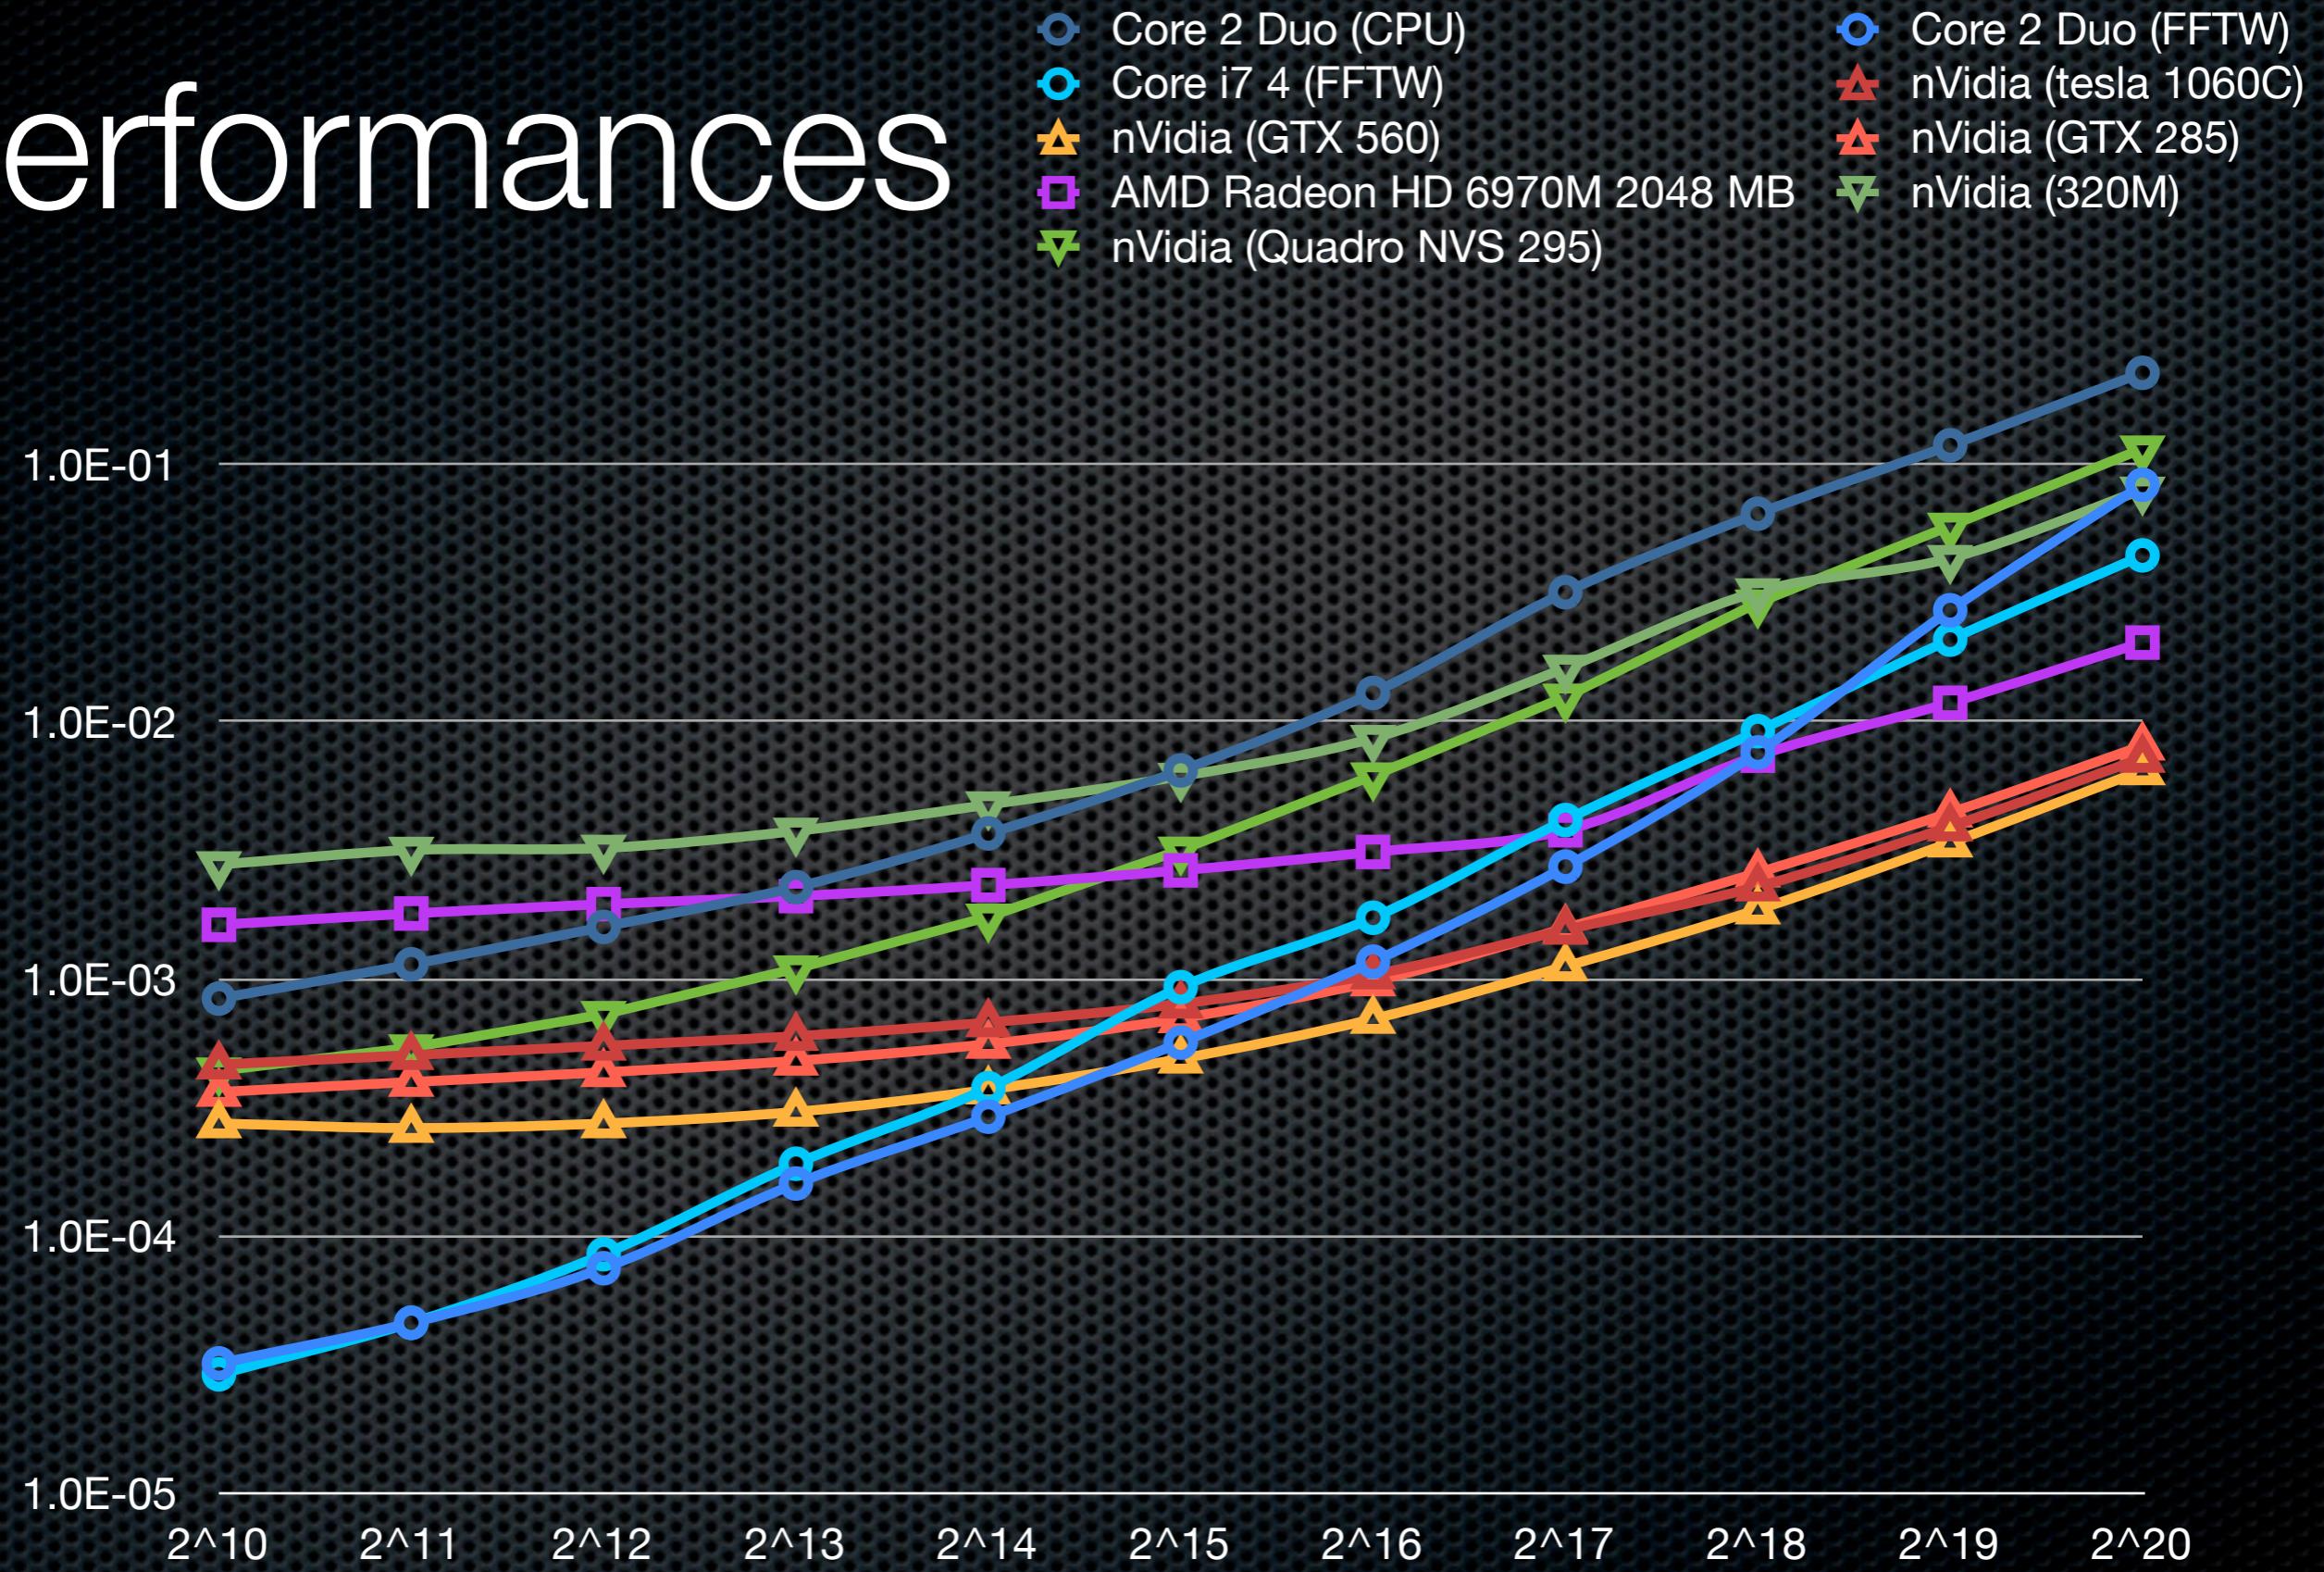
- From Temporal space to Spectral
- Discrete transformation
- Based on power of prime number
- Divide & Conquer algorithm
- $N \log(N)$ complexity



Kernel & FFTW

- From the creator of the official OpenCL FFT library
- Release by Apple under MIT type license
- Has similar performances to the CUDA FFT library
- FFTW is a library written for CPU
- Used by Matlab in fft
- Suppose to be the best library for FFT on CPU
- GPL license

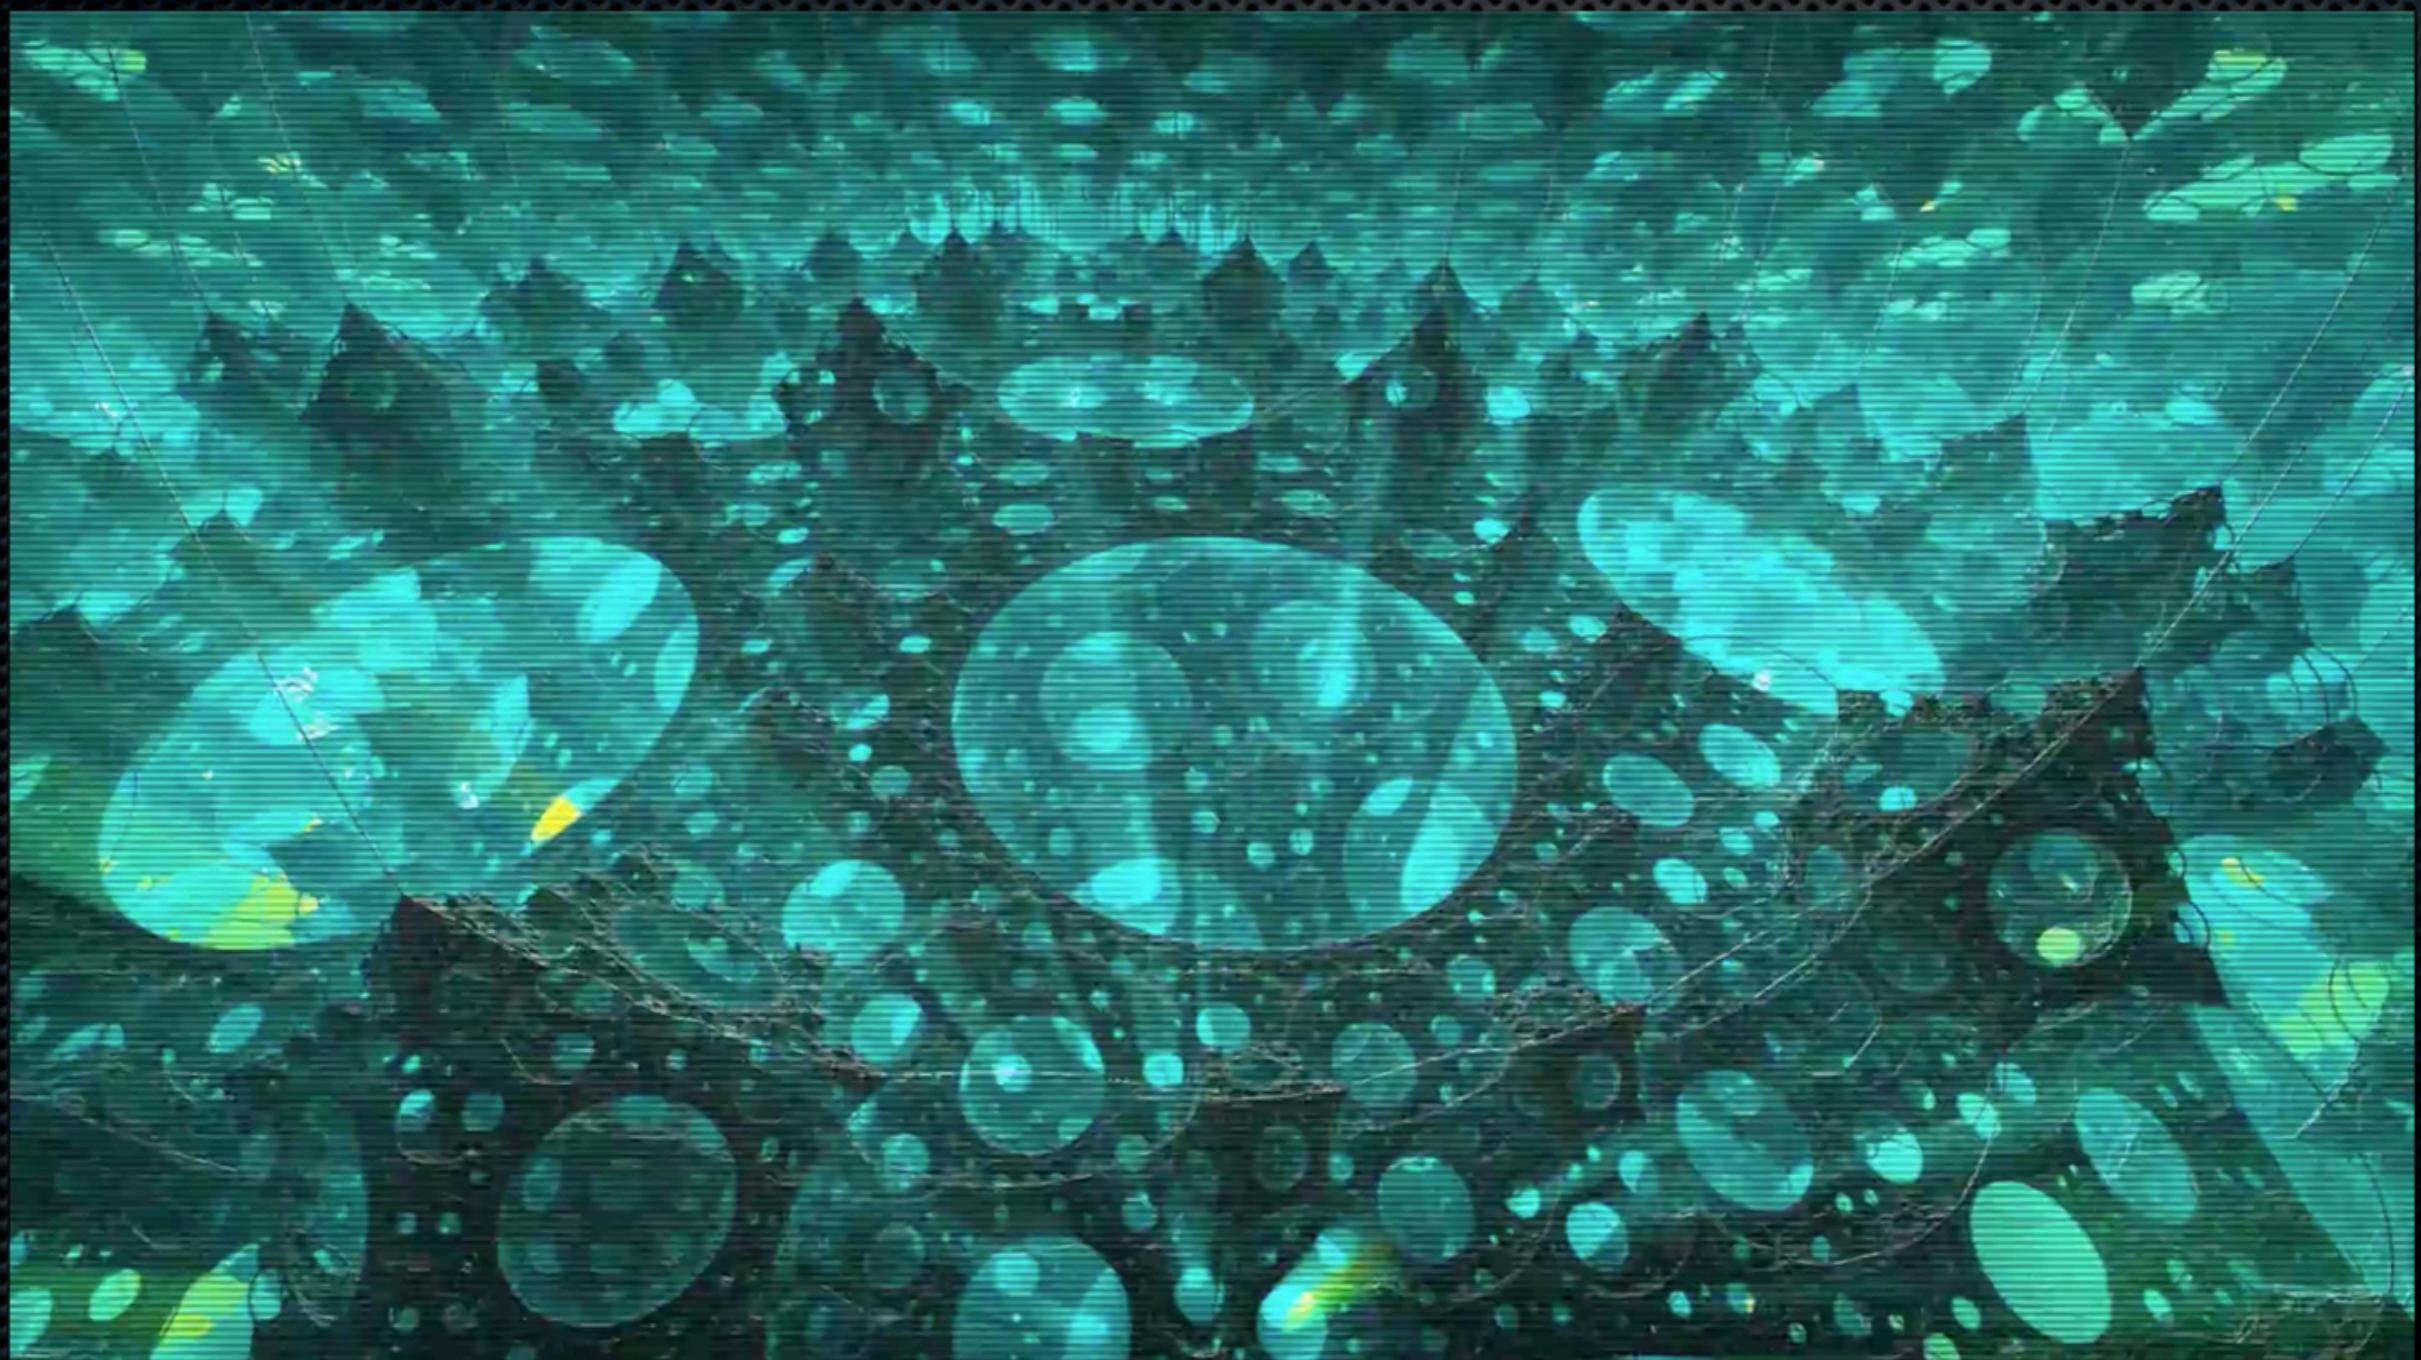
Performances



Conclusion

- OpenCL is a mature technology
- Performances are at least at the level of CUDA
- Ability to use it on CPU is a plus
- GPGPU offers excellent performances at a reduce cost
 - Cheaper than Hardware
 - Shorter development time

Questions?



ray marching using OpenCL+OpenGL - Hartverdrahtet - Akronyme Analogiker | 4k Revision 2012